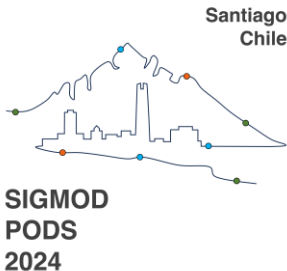


# Can Learned Indexes be Built Efficiently?

## A Deep Dive into Sampling Trade-offs

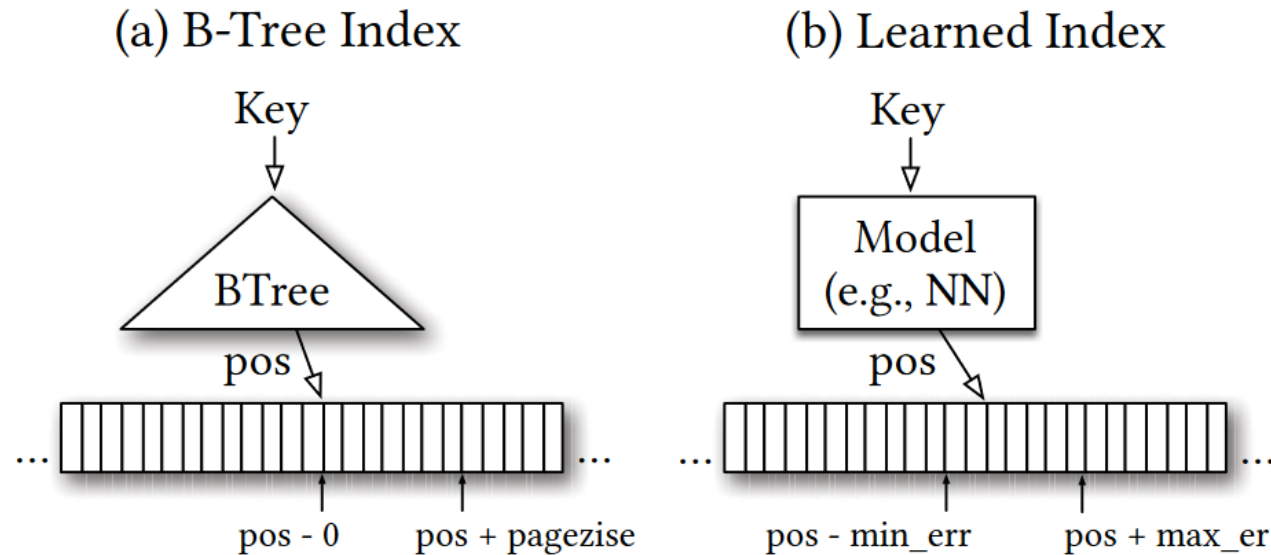
Minguk Choi, Seehwan Yoo, Jongmoo Choi

Dankook University, South Korea



# 1. Introduction

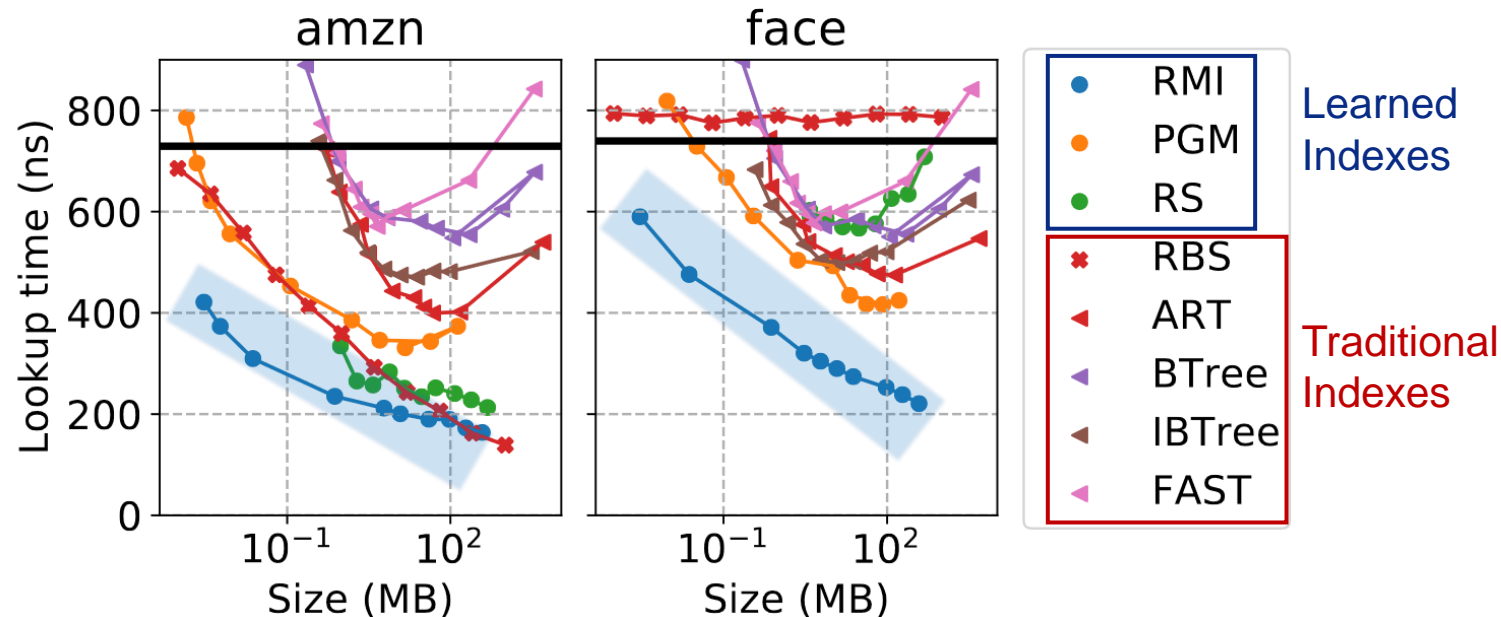
- Learned index
  - Index structure employs machine learning techniques
  - View the index as a model that **predicts the position of a key** in sorted array



(The Case for Learned Index Structures, SIGMOD '18)

# 1. Introduction

- Learned index
  - **Space-efficient** by effectively compressing data distribution through model
  - Pareto optimal in terms of index size and lookup latency in read-only workload
    - No alternative exists that has both a smaller size and lower latency

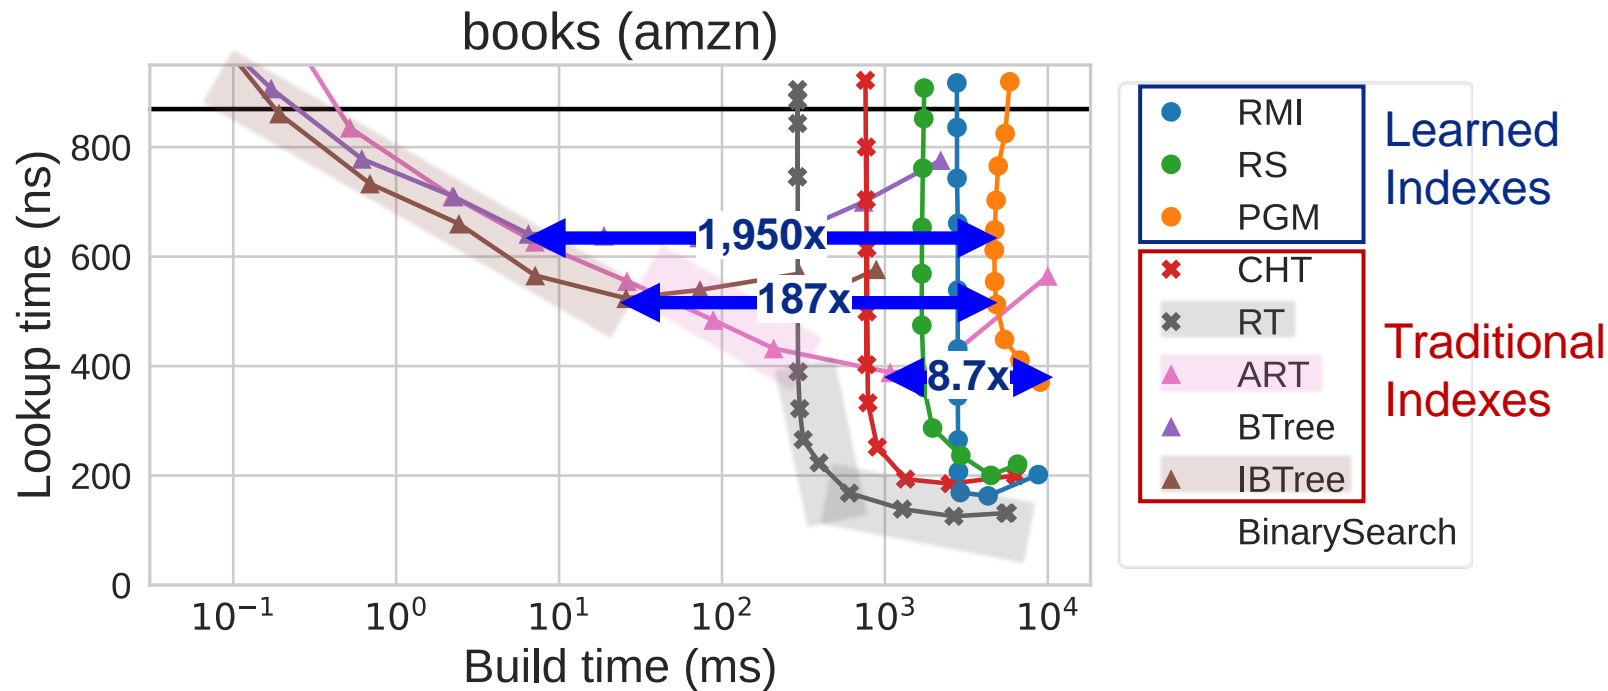


(Benchmarking Learned Indexes, VLDB '20)

# 1. Introduction

- Limitation of learned index : **Long index build time**

- Significantly (up to about 2,000x) longer than traditional indexes
  - Not Pareto optimal (build-efficient) for build time and lookup latency
- Still, there are **application** (e.g., LSM-Tree) where the index build time is crucial



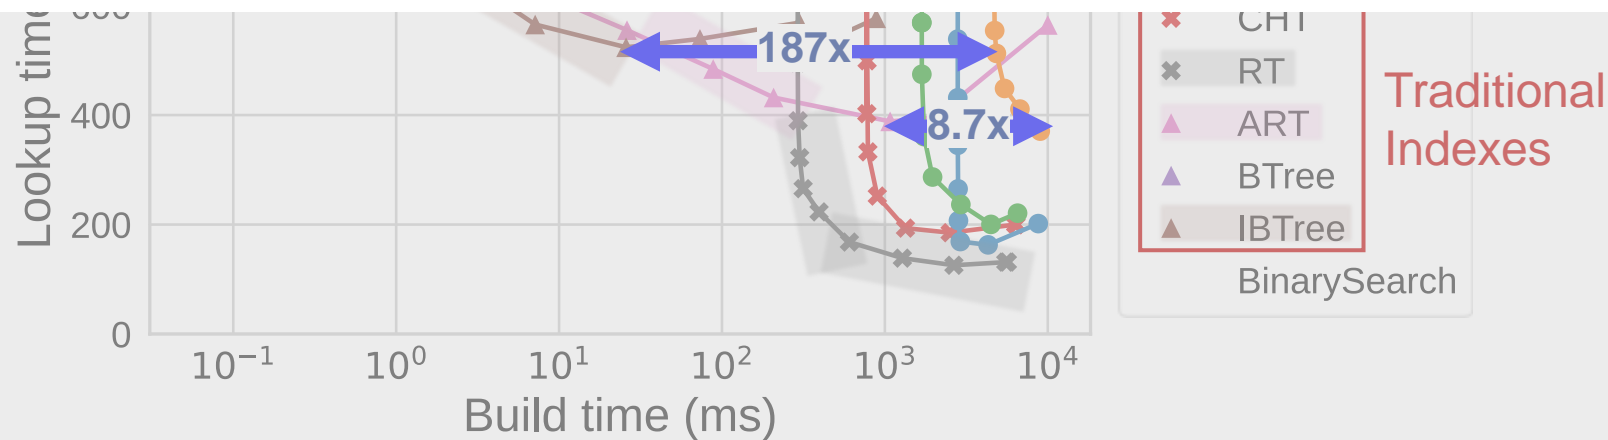
# 1. Introduction

- Limitation of learned index : **Long index build time**
  - Significantly (up to 1,950x) longer than traditional indexes



Long build time has been identified as **a high priority** for future work in various papers

RMI (SIGMOD `18), RadixSpline (aiDM `20), PGM-Index (VLDB `20), SOSD (VLDB `20), Critical-RMI (VLDB `22)



# 1. Introduction

- The primary reason for long build time of learned index

$$\text{Index build time} = \text{Per - element overhead} \times \text{Number of elements}$$

- 1) Higher per-element overhead
- 2) Complete traversal and training

- To mitigate per-element overhead
  - Light-weight model : RadixSpline (aiDM `20), Bourbon (OSDI `20)
    - It still shows longer build time than traditional indexes

This study began with the question 🤔...

Since a learned index uses a model,

**Can't it learn efficiently even  
with less data?**

# 1. Introduction

- Our Approach : **Sampling**

- While sampling may seem simple and even naïve, it is indeed complex


- Challenges

1. **Absence** of benchmark for sampling applied indexes
2. **Losing** error-bound property due to sampling loss
3. **Complex** trade-offs in terms of model, index, and micro-architecture




## 2. Background

- Workload : Read-only in-memory
  - Practical beginning point of learned index
  - Dataset ( $D$ ) : Sorted array of unique integer keys without duplicates
  - Lookup : Find the position of a lookup key  $k$  in  $D$ 
    - ① Prediction : Estimate the position of  $k$  as  $p$
    - ② Correction : Find the exact position of  $k$  based on  $p$




## 2. Background

- Workload : Read-only in-memory
  - Error-bound property
    - $\forall k \in D, Error(k) = |Pred(k) - Pos(k)| \leq \varepsilon$  (*= error - bound*)
    - $k$  exist in correction range ( $= [p - \varepsilon, p + \varepsilon]$ )  $\rightarrow$  binary search
  - Important for robustness,
    - Especially where correction is expensive
    - E.g., Disk or remote I/O environments



# 3. Design



## 1. Unified sampling algorithm & implementation

 **BASIL** (Benchmark of Sampling Applied Learned Indexes)

### 1) Unified sampling algorithm

- **Systematic** sampling : extract every  $I^{th}$  ( $I$  = sampling interval) key from 0-th key to the last key
  - ✓ Pros : Simple, universal, no decision/reordering cost
  - ✓ Cons: Not optimal (other methods, e.g., adaptive, should be explored)

### 2) Unified sampling implementation

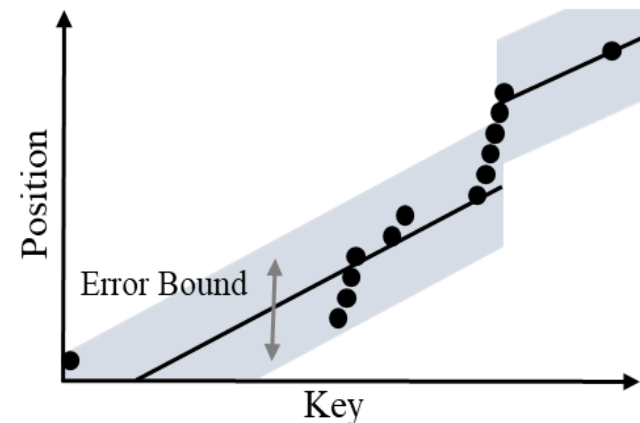
- All indexes access and train only sample key-value data from entire dataset

# 3. Design


## 2. Sample learning algorithm

- EB-PLA (Error-bounded Piece-wise Linear Approximation) model
  - Without sampling, model trains **all** keys with error-bound  $\varepsilon \rightarrow \text{Error}(k) \leq \varepsilon$
  - With sampling, model trains only **sample**  $I^{th}$  keys with error-bound  $\varepsilon \rightarrow \text{Error}(k) \leq \varepsilon$ 
    - **Losing** of the error-bound property, which is learning objective of the model

(a) Train All ( $I = 1$ )  
with  $\varepsilon = 3$



(b) Train Sample ( $I = 3$ )  
with  $\varepsilon = 3$




# 3. Design

## 2. Sample learning algorithm


- Sample EB-PLA algorithm

- Refine error-bound due to sampling loss ( $\varepsilon' = \varepsilon + I - 1$ )
  - In Fig. (c), preserves error-bound property, but cannot guarantee desired error-bound ( $\varepsilon$ )
- Replace the learning error-bound to  $\delta (= \varepsilon - I + 1)$  for desired error-bound ( $\varepsilon$ )
  - In Fig. (d), preserves error-bound ( $\varepsilon$ ) by learning less data with smaller & stricter error bound


(a) Train All ( $I = 1$ )  
with  $\varepsilon = 3$




(b) Train Sample ( $I = 3$ )  
with  $\varepsilon = 3$



(c) Refine Error-bound  
 $\varepsilon$  from 3 to 5



(d) Train Sample ( $I = 3$ )  
with  $\delta = 1$

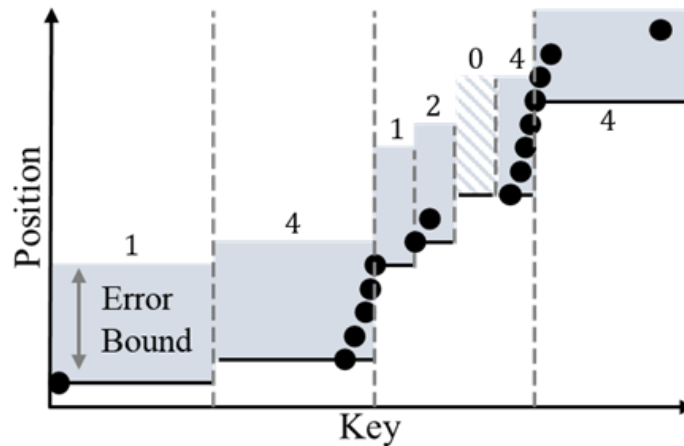


# 3. Design

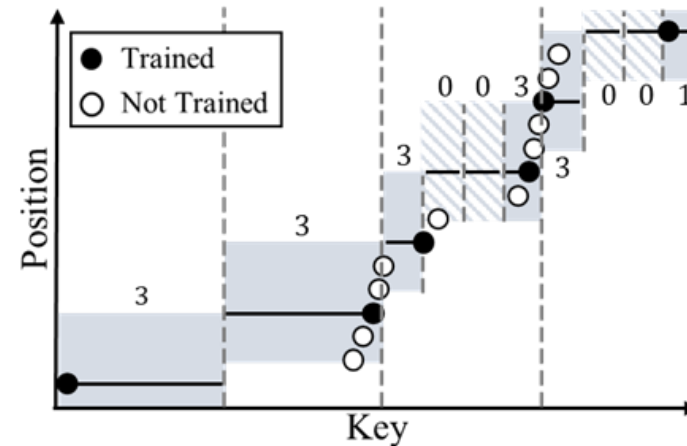
## 2. Sample learning algorithm

- Sample EB-Histogram algorithm

- Fig (a), Train all with error-bound  $\varepsilon \rightarrow \forall k \in D, k \in [p, p + \varepsilon]$
- Fig (b), Train sample with error-bound  $\delta (= \varepsilon - I + 1) \rightarrow \forall k \in D, k \in [p - I + 1, p + \delta]$ 
  - Preserve Correction length ( $\varepsilon + 1 = \delta + I$ )



(a) EB-Histogram ( $I = 1, \varepsilon = 5$ )

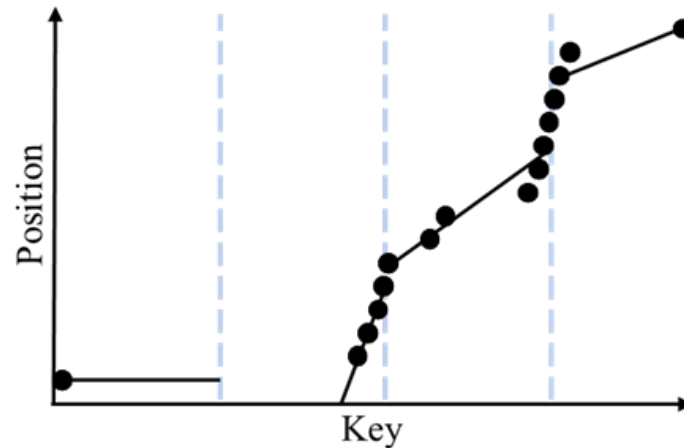


(b) Sample EB-Histogram ( $I = 1, \delta = 3, \varepsilon = 5$ )


# 3. Design

## 2. Sample learning algorithm

- PLR with Simple Linear Regression
  - Fig. (a), Model itself cannot guarantee an error-bound regardless of sampling
    - To guarantee error-bound, Measuring all data errors after training causes significant overhead.
  - Fig. (b), Sampling can decrease accuracy(MSE), but error-bound property doesn't change.



(a) Simple Linear Regression ( $I = 1$ )



(b) Simple Linear Regression ( $I = 3$ )


# 3. Design

## 3. Internal changes due to sampling

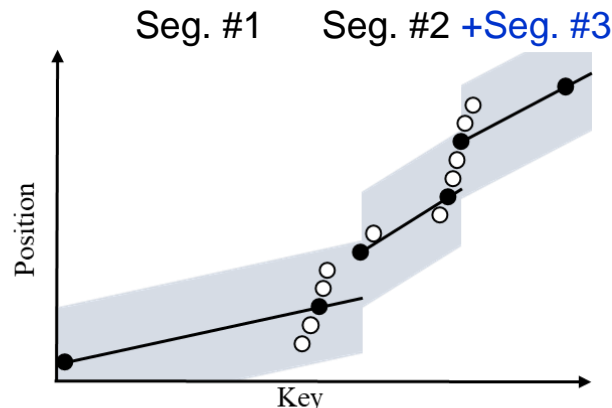
- Depend on segmentation manner

### 1) **Dynamic** segmentation (EB-PLA, EB-Histogram)


- Definition : Dynamically segment key ranges according to the distribution
- Trade-off : Decrease build time but aggressive sampling can **increase # of segments (bins)**




(a) EB-PLA ( $I = 1, \varepsilon = 3$ )



(b) Sample EB-PLA ( $I = 1, \delta = 1, \varepsilon = 3$ )



(a) EB-Histogram ( $I = 1, \varepsilon = 5$ )



(b) Sample EB-Histogram ( $I = 1, \delta = 3, \varepsilon = 5$ )



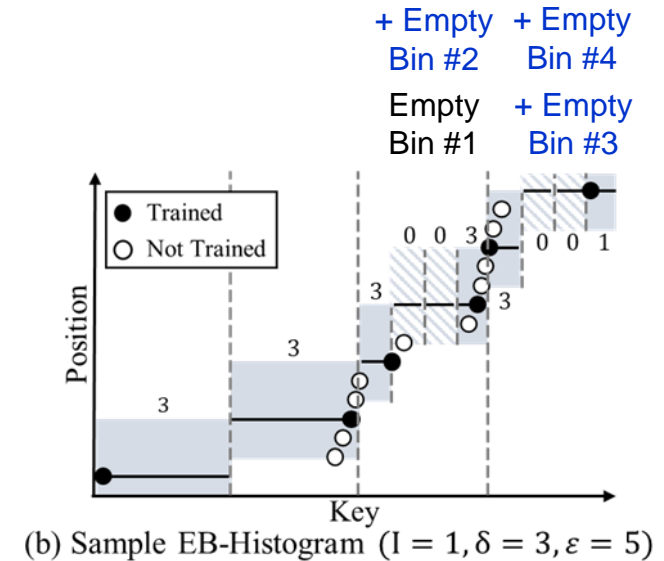
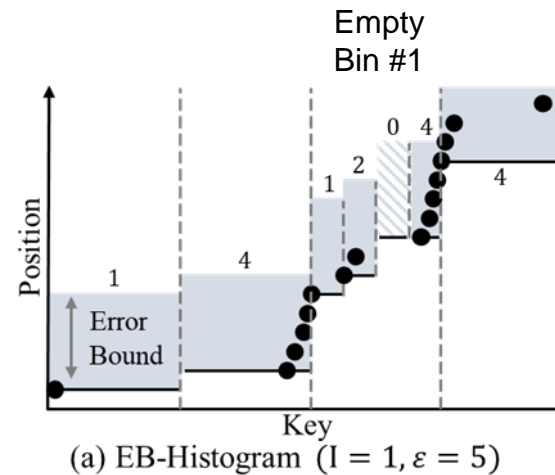
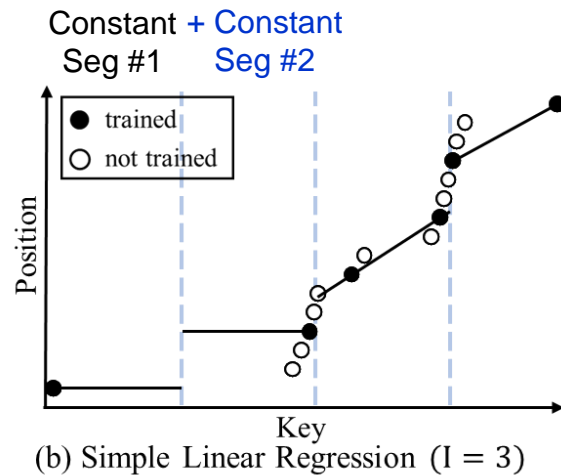
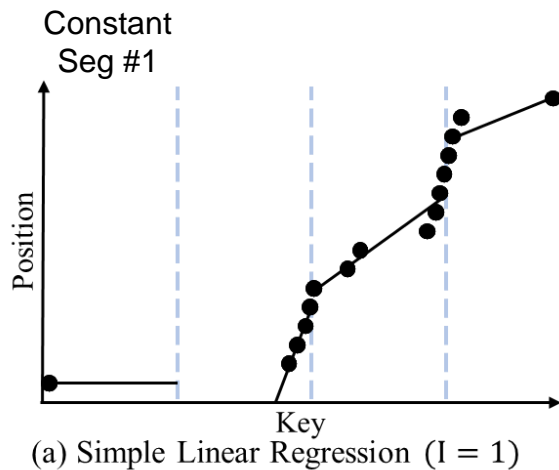
# 3. Design

## 3. Internal changes due to sampling

- Depend on segmentation manner

### 2) Fixed segmentation (PLR, EB-Histogram)

- Definition : Segment key ranges into a fixed number of segments
- Trade-off : Decrease build time but aggressive sampling can increase # of underfitting segments



# 4. Evaluation Setup



## BASIL (Benchmark of Sampling Applied Learned Indexes)

- Applied sampling to 7 indexes, prefixed with “s”
  - 3 Learned, 2 Histogram, 3 Tree-based indexes


Type	Index	Internal Model	Correction Search
Learned	sRMI	Simple Linear Regression	Exponential Search
Learned	sPGM / sRS	Sample EB-PLA	Binary Search
Histogram	sCHT	Sample EB-Histogram (Equal-width)	
Histogram	sRT	Sample Histogram (Equal-width)	
Tree-based	sART / sB+-Tree/ sIB-Tree	-	

## 4. Evaluation Setup



### BASIL (Benchmark of Sampling Applied Learned Indexes)

- Datasets : 6 representative datasets with 200 million key-value pairs
- Workload : Lookup uniform random 10 million keys from the dataset.
- Environment: Intel(R) Xeon(R) Gold 6338 CPU 2.00 GHz, 48 MB L3 with 512 GB of main memory



# 5. Evaluation

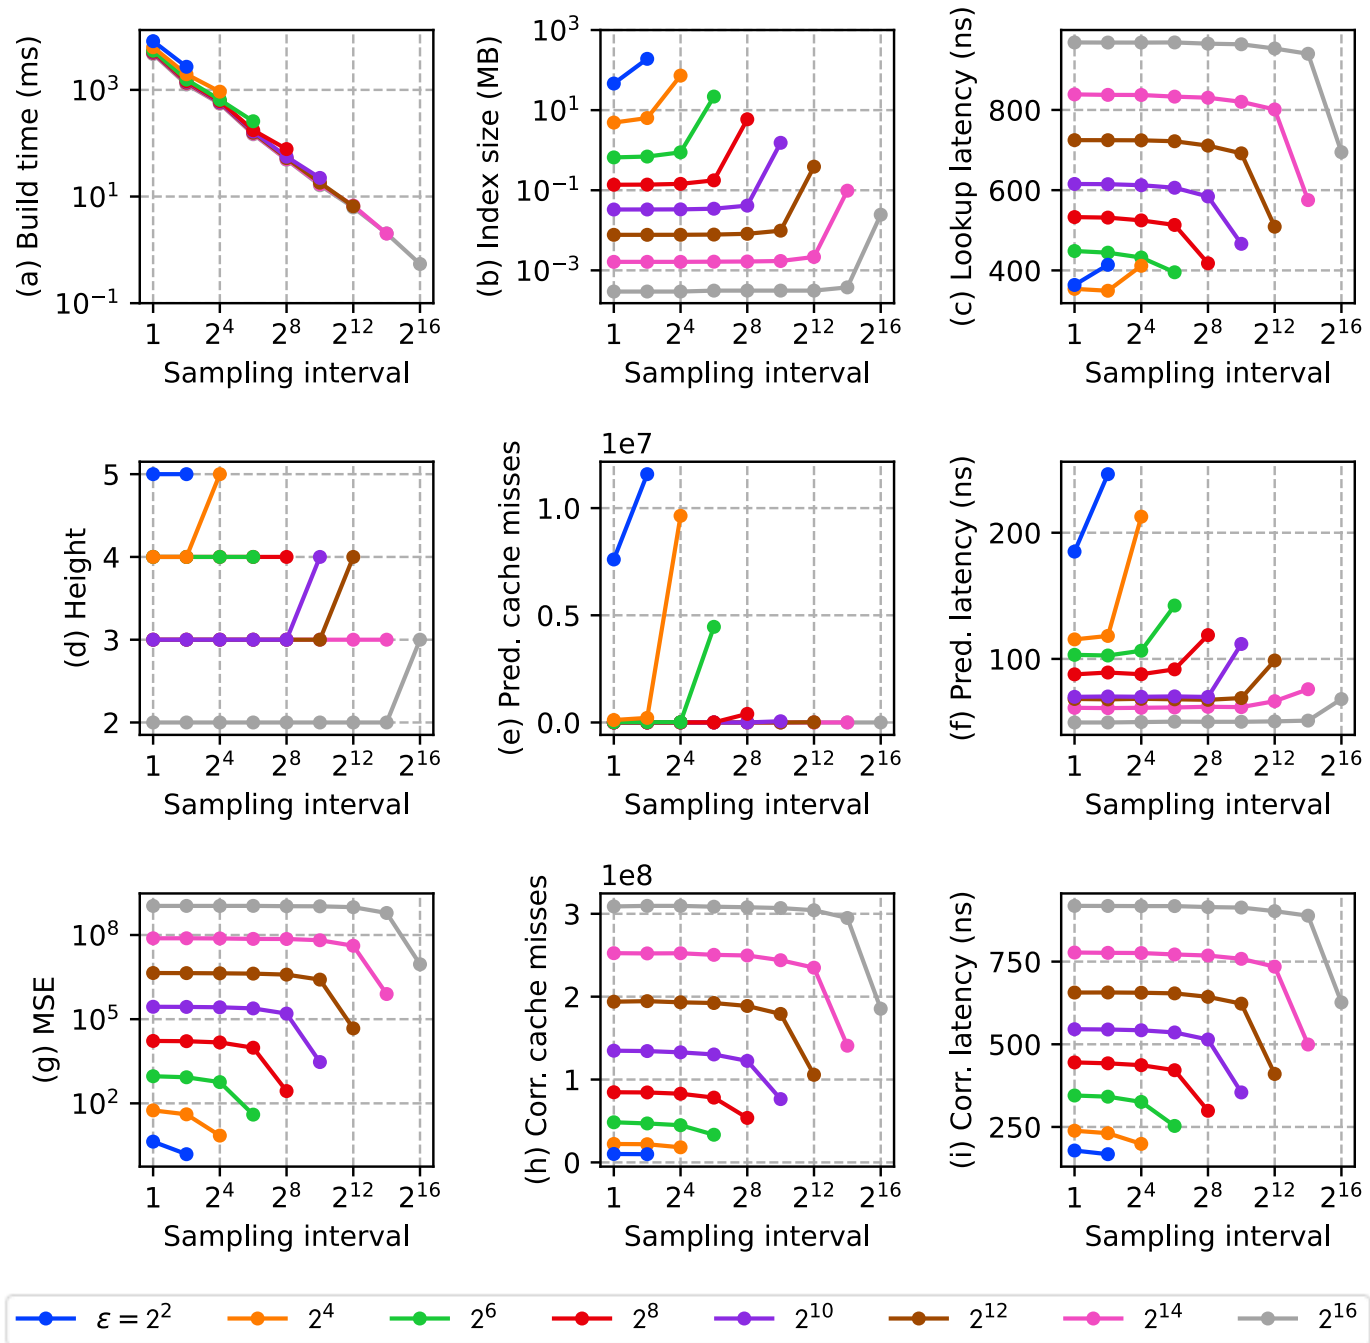
## 1. Sampling Trade-offs

### ■ Index

- sPGM with Sample EB-PLA

### ■ Metrics

- **Index:** (a) Build Time, (b) Size, (c) Latency, (f) Pred. latency, (i) Corr. latency
- **Model:** (d) Height, (g) MSE (Accuracy)
- **Micro-architecture:** (e) Pred. Cache Miss, (f) Corr. Cache Miss

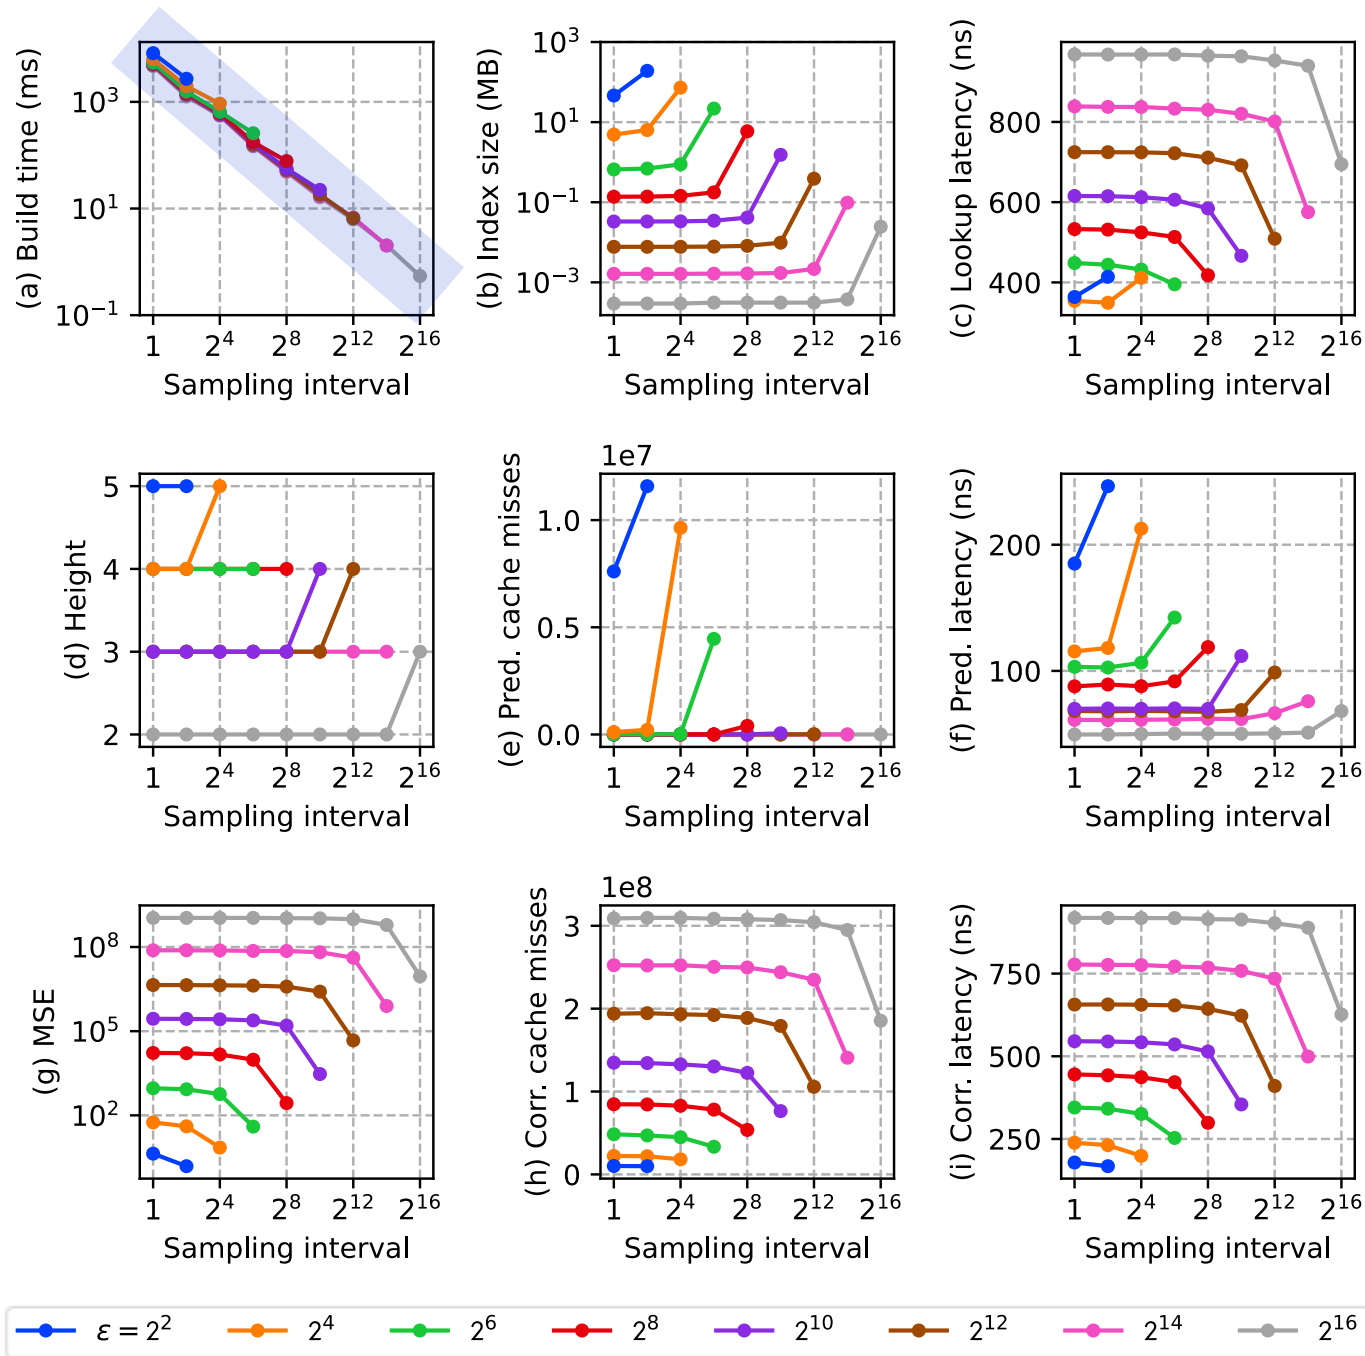


# 5. Evaluation

## 1. Sampling Trade-offs

- When sampling interval (I) increases, **(a) build time decreases** by order of magnitude

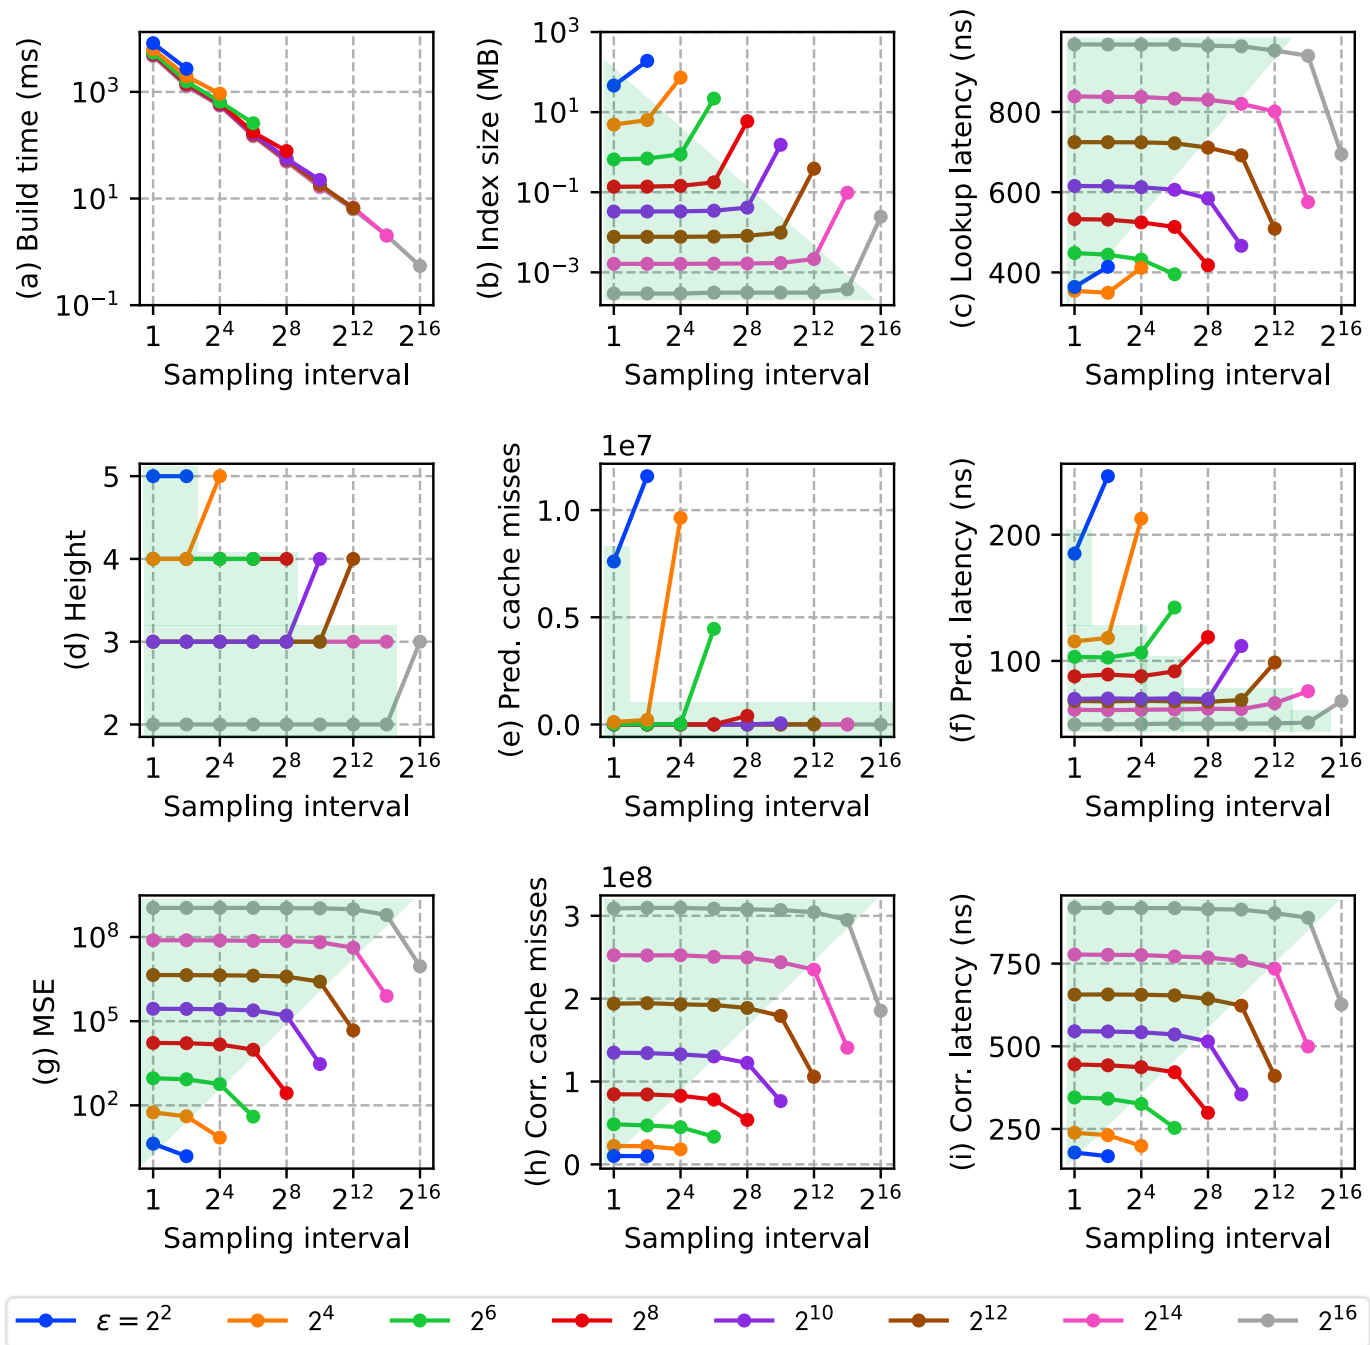
Dataset : History, Error bound ( $\varepsilon \in [2^2, 2^{16}]$ ), Sampling interval ( $I \in [2^0, \varepsilon (\leq 2^{16})]$ )



# 5. Evaluation

## 1. Sampling Trade-offs

- Each error-bound has **threshold interval ( $I_{TH}$ )**
  - mostly  $\varepsilon = I^{TH}$
- Until** threshold interval ( $I_{TH}$ ), (b-i) the rest of metrics **remain consistent**



# 5. Evaluation

## 1. Sampling Trade-offs

- After threshold interval  $I_{TH}$ ,

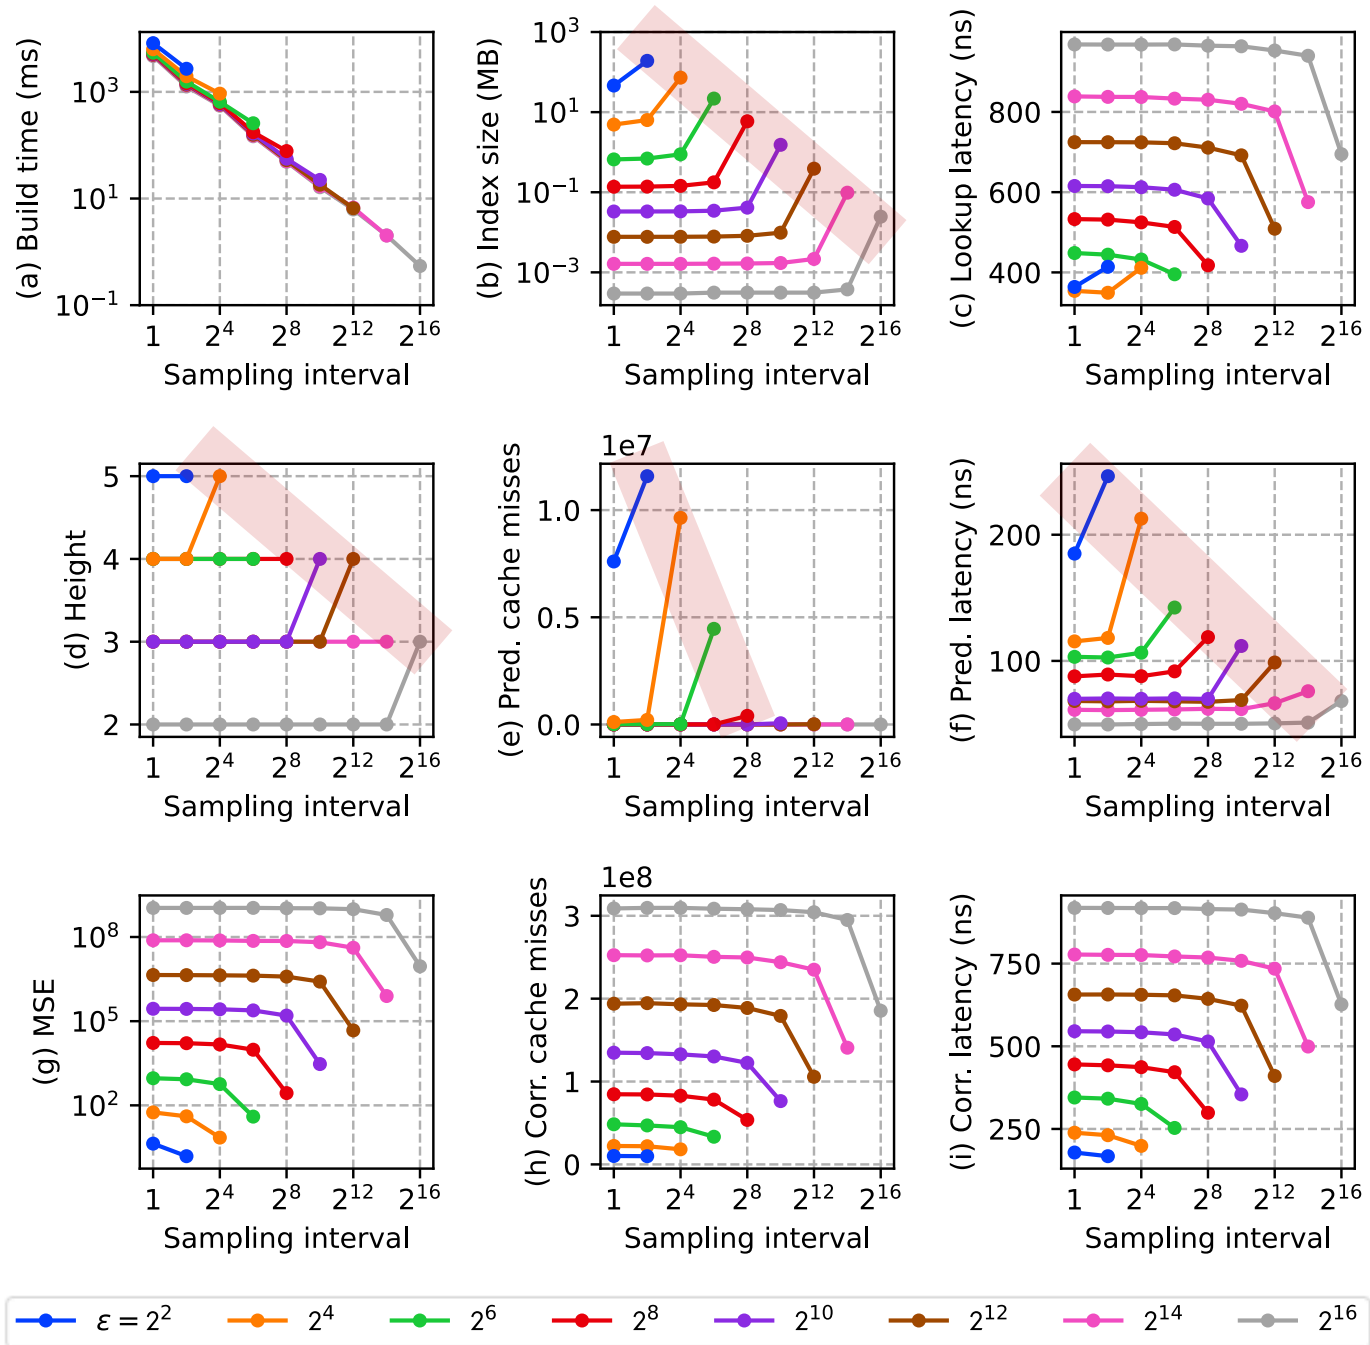
# of linear segments  $\uparrow$

→ (b) Size  $\uparrow$

(d) Height  $\uparrow$

→ (e) Pred. cache miss  $\uparrow$ ,

(f) Pred. latency  $\uparrow$



# 5. Evaluation

## 1. Sampling Trade-offs

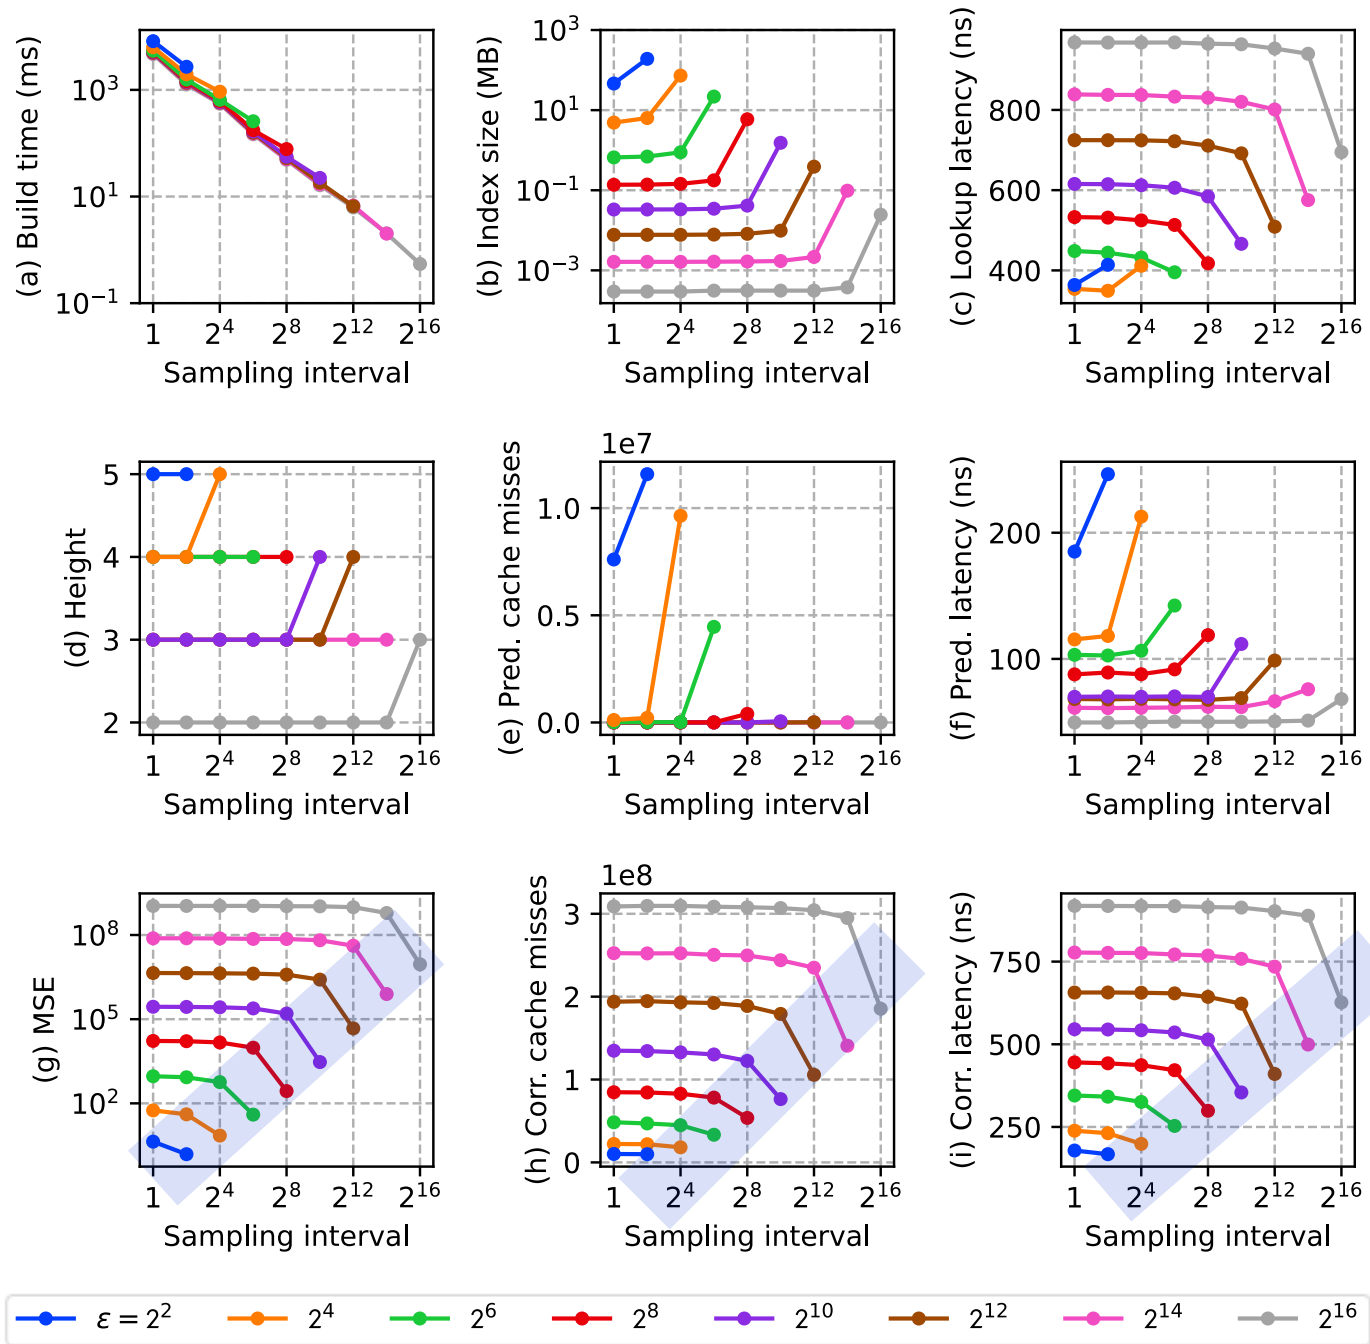
- After threshold interval  $I_{TH}$ ,

# of linear segments  $\uparrow$

→ (g) MSE  $\downarrow$

→ (h) Corr. cache miss  $\downarrow$

(i) Corr. latency  $\downarrow$

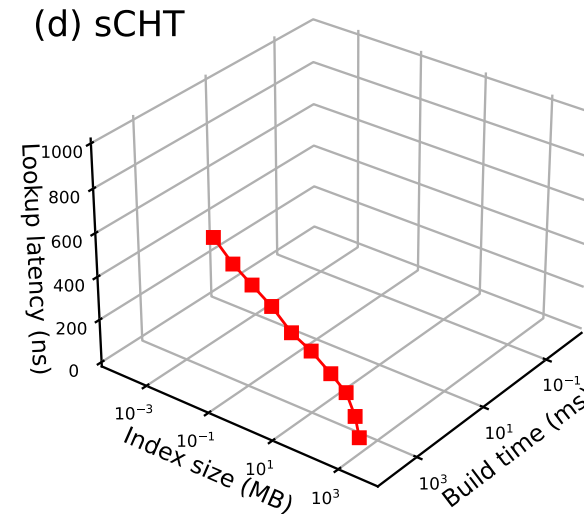
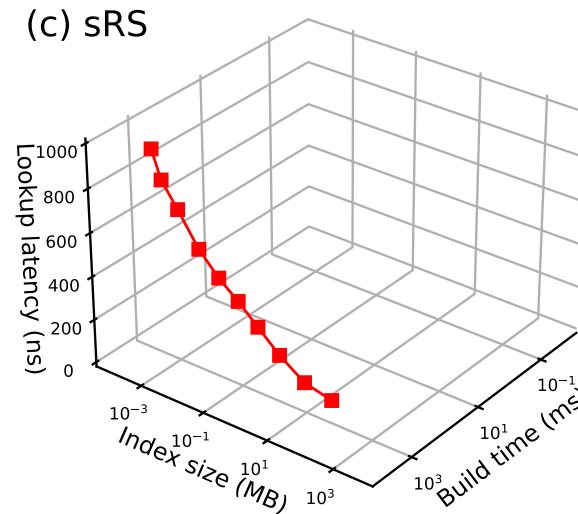
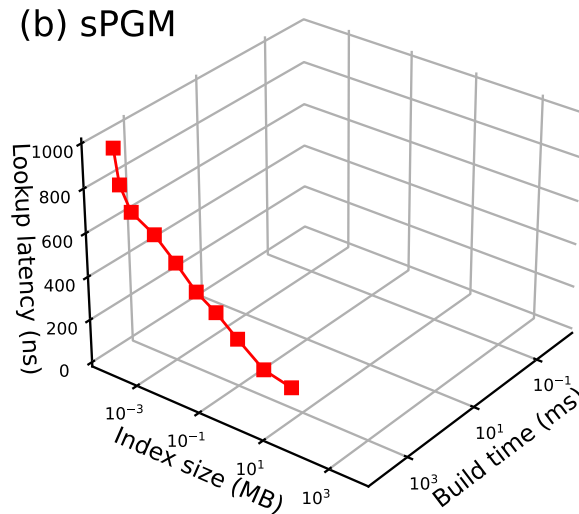
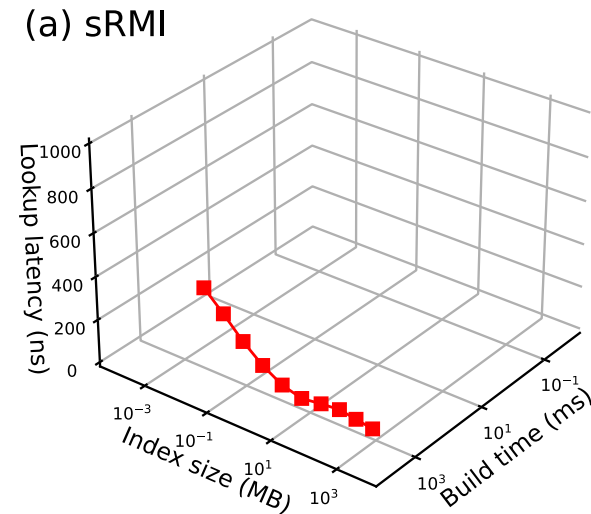




# 5. Evaluation

## 2. Design Space of Learned Indexes

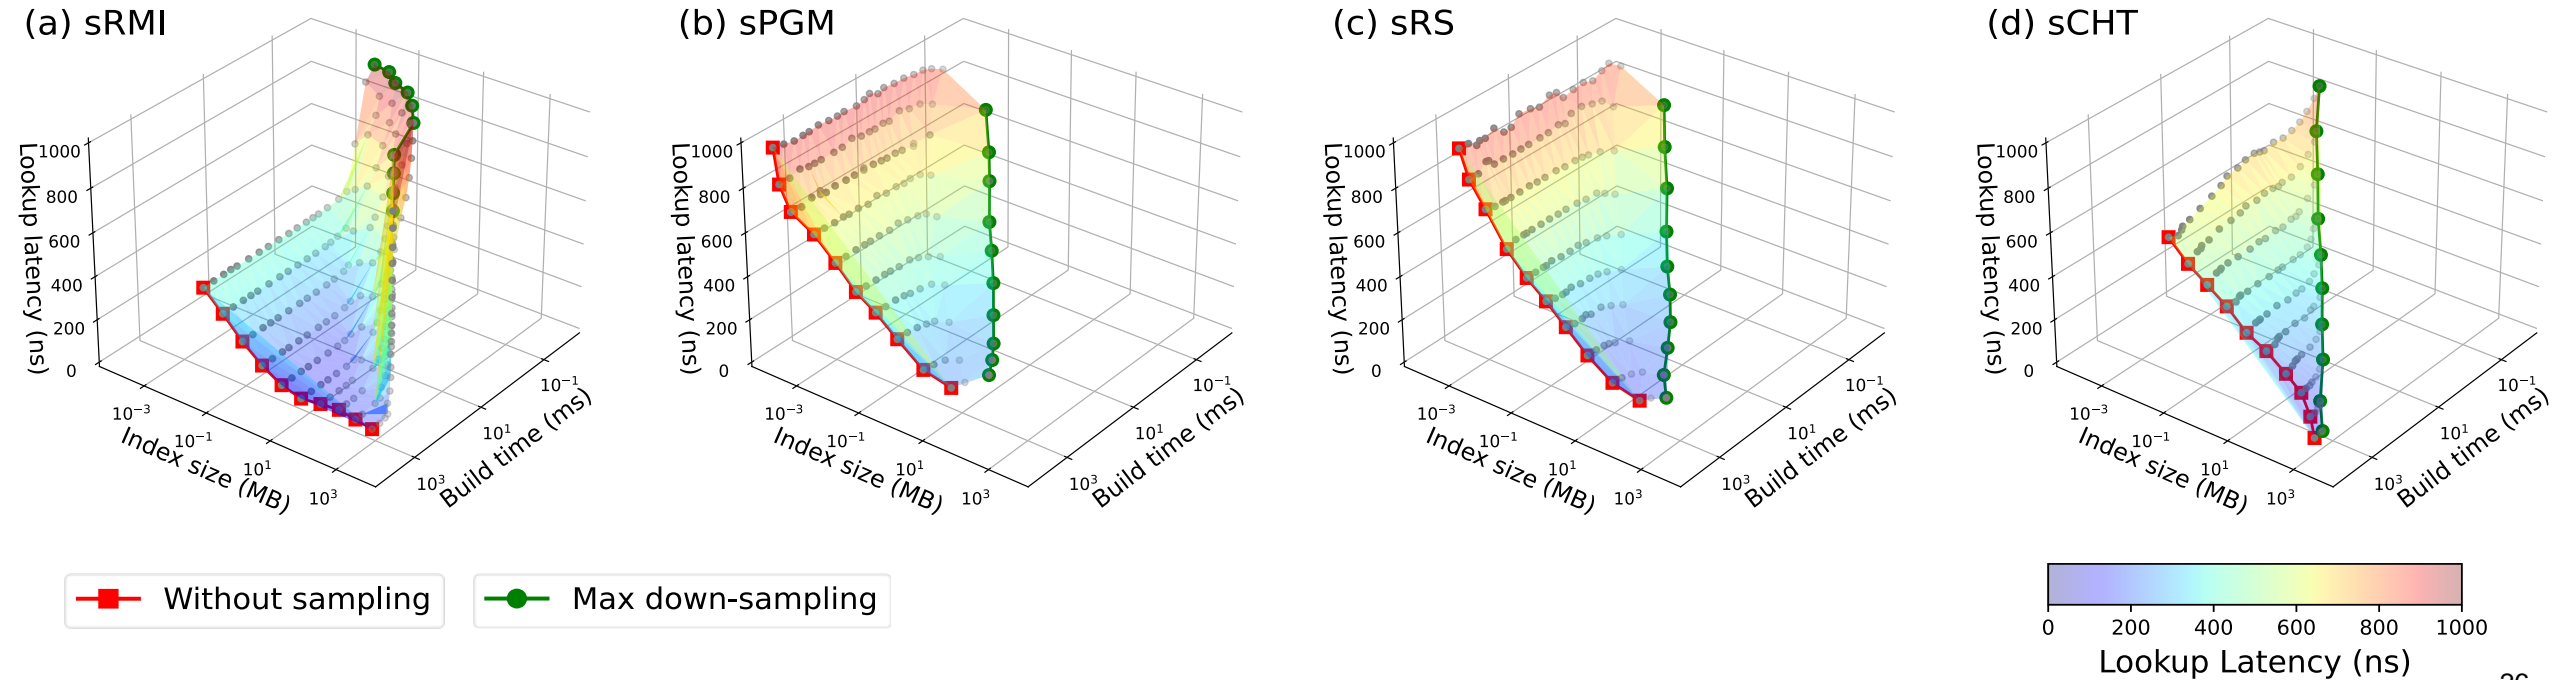
- **Absence** of trade-offs between build time, index size, and lookup latency
  - Incur significant build times regardless of size and lookup



—■— Without sampling

# 5. Evaluation

- 2. Design Space of Learned Indexes
  - Sampling **introduces** trade-offs between build-time, size, and lookup latency
    - Broaden** design space of learned indexes from 2D to 3D

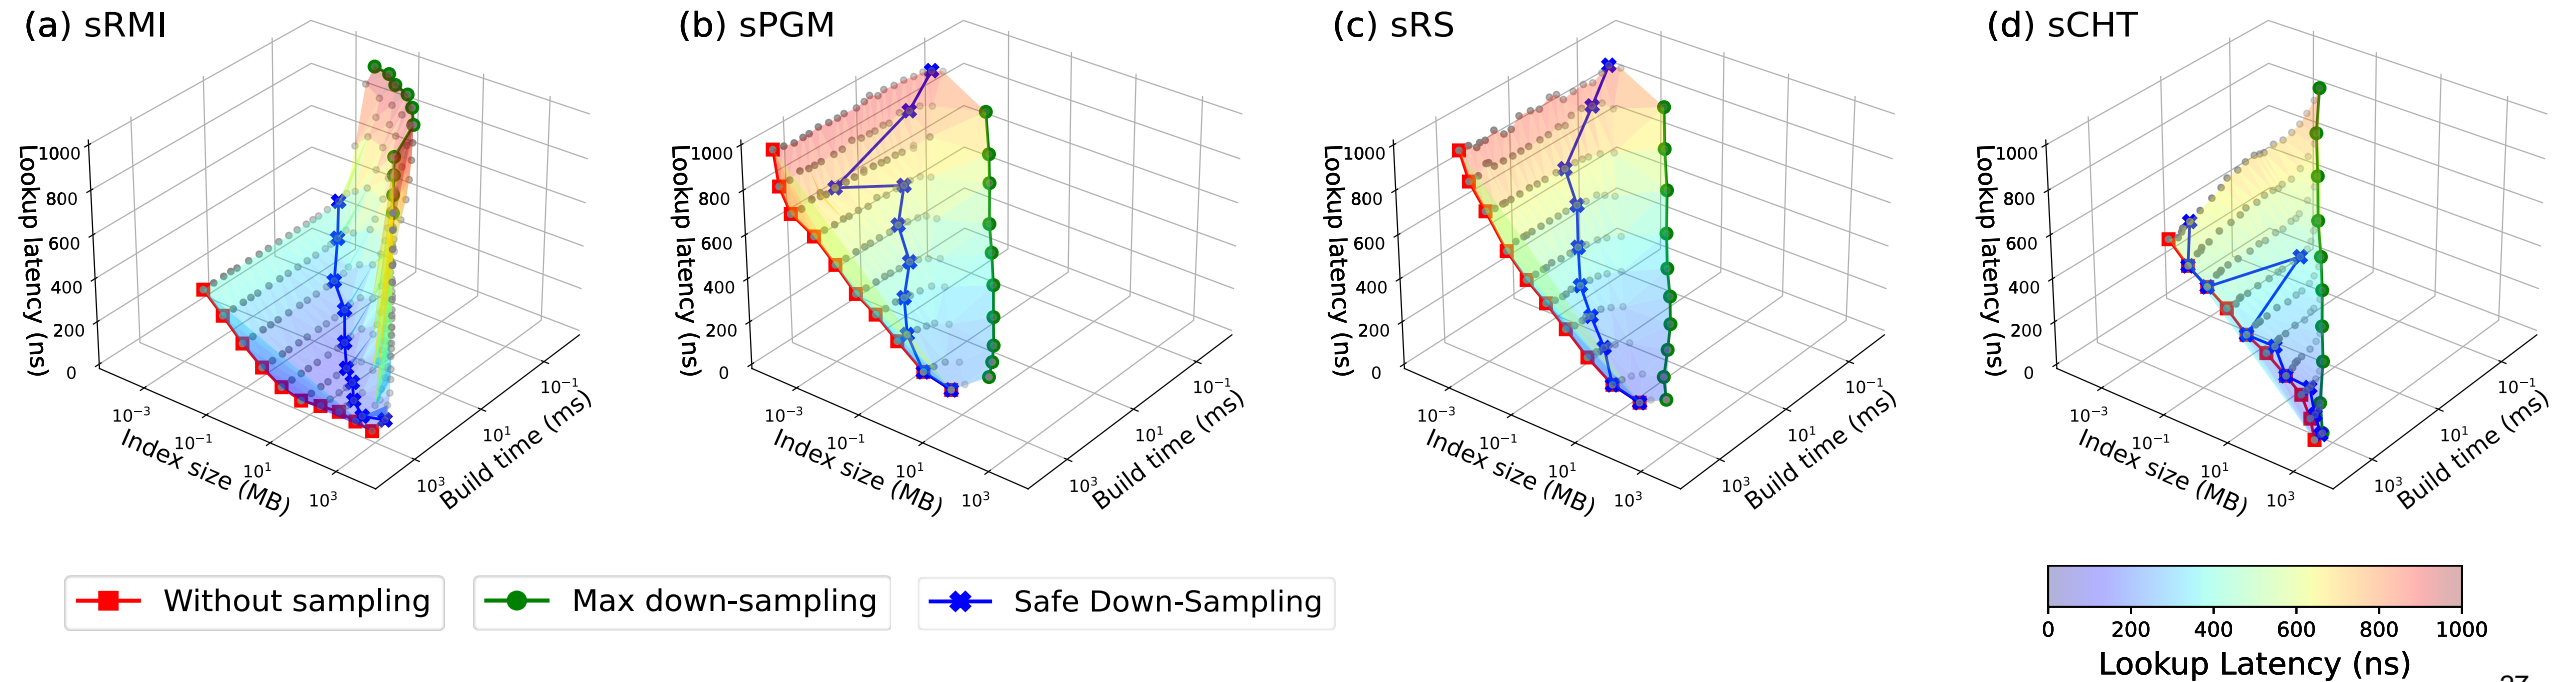


# 5. Evaluation

## ■ 3. Build Speed-up

- Question. How much can sampling reduce build time without significantly degrading index performance?

➤ **Safe** Down-sampling where size & lookup latency increase by **less than 5%**




# 5. Evaluation

## 3. Build Speed-up

- Question. How much can sampling reduce build time without significantly degrading index performance?

➤ **Safe** Down-sampling where size & lookup latency increase by **less than 5%**



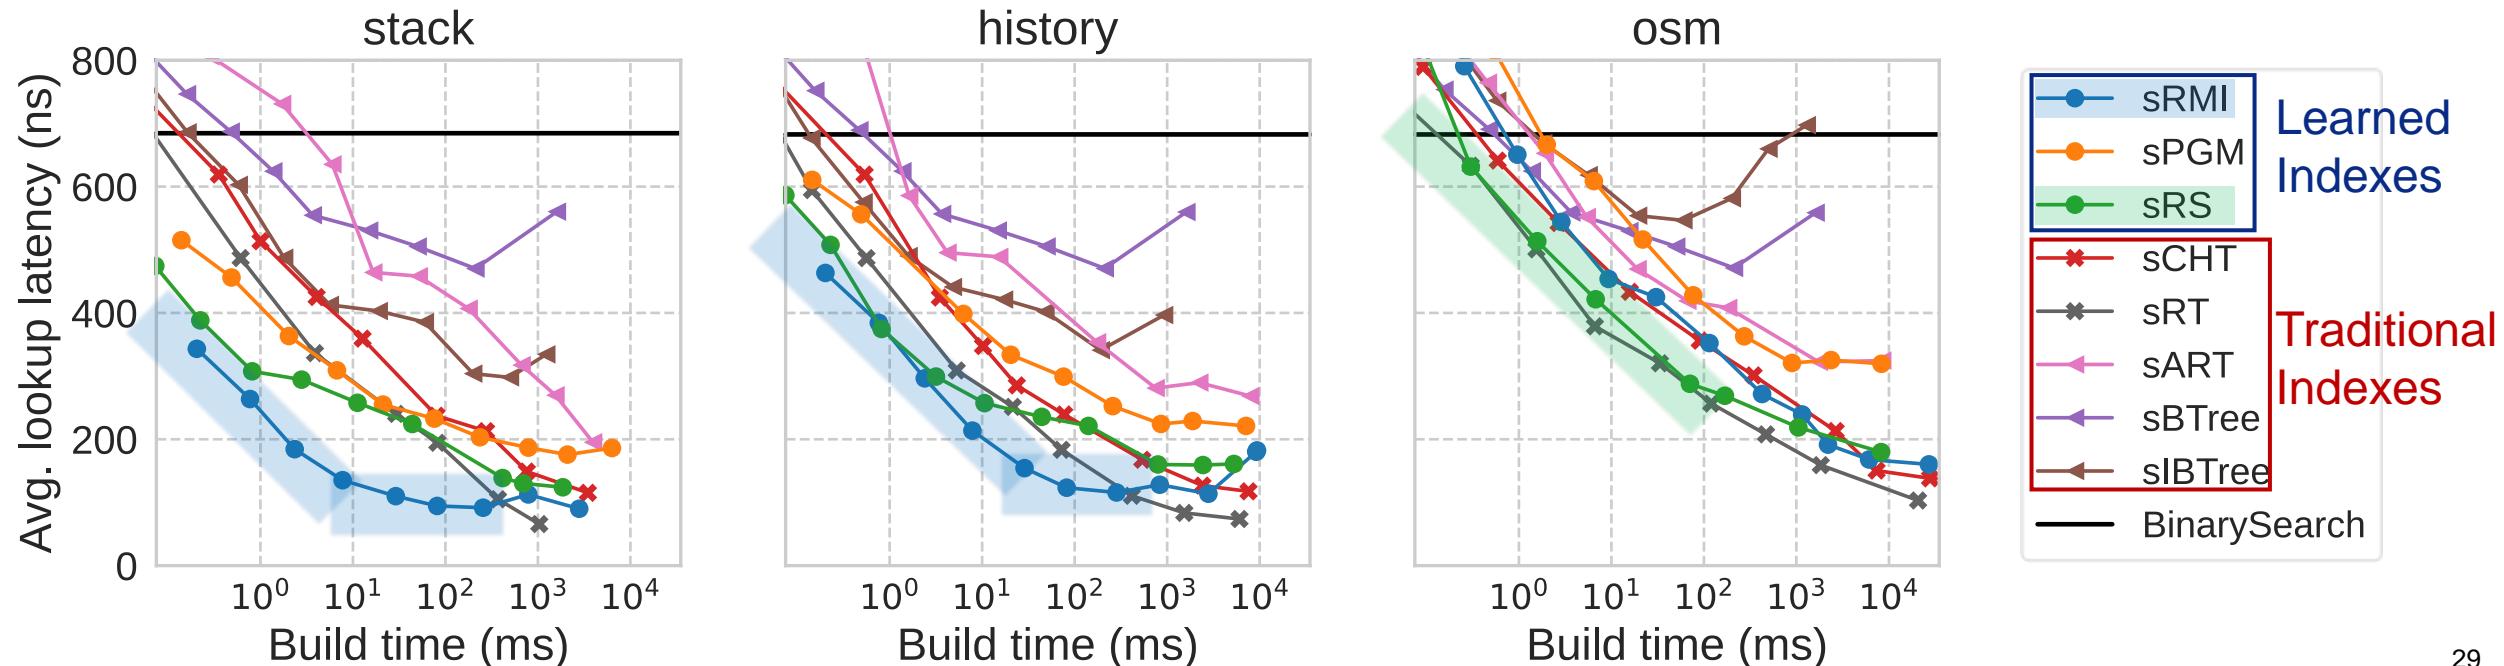
# 5. Evaluation

## ■ 4. Pareto Optimal Analysis

- Question. Can learned indexes be built efficiently (in terms of build time and lookup latency than traditional indexes through sampling)?

- **Pareto optimal** (build-efficient) in terms of build time and **average** lookup latency

➤ no alternative that has both shorter build time and lower latency



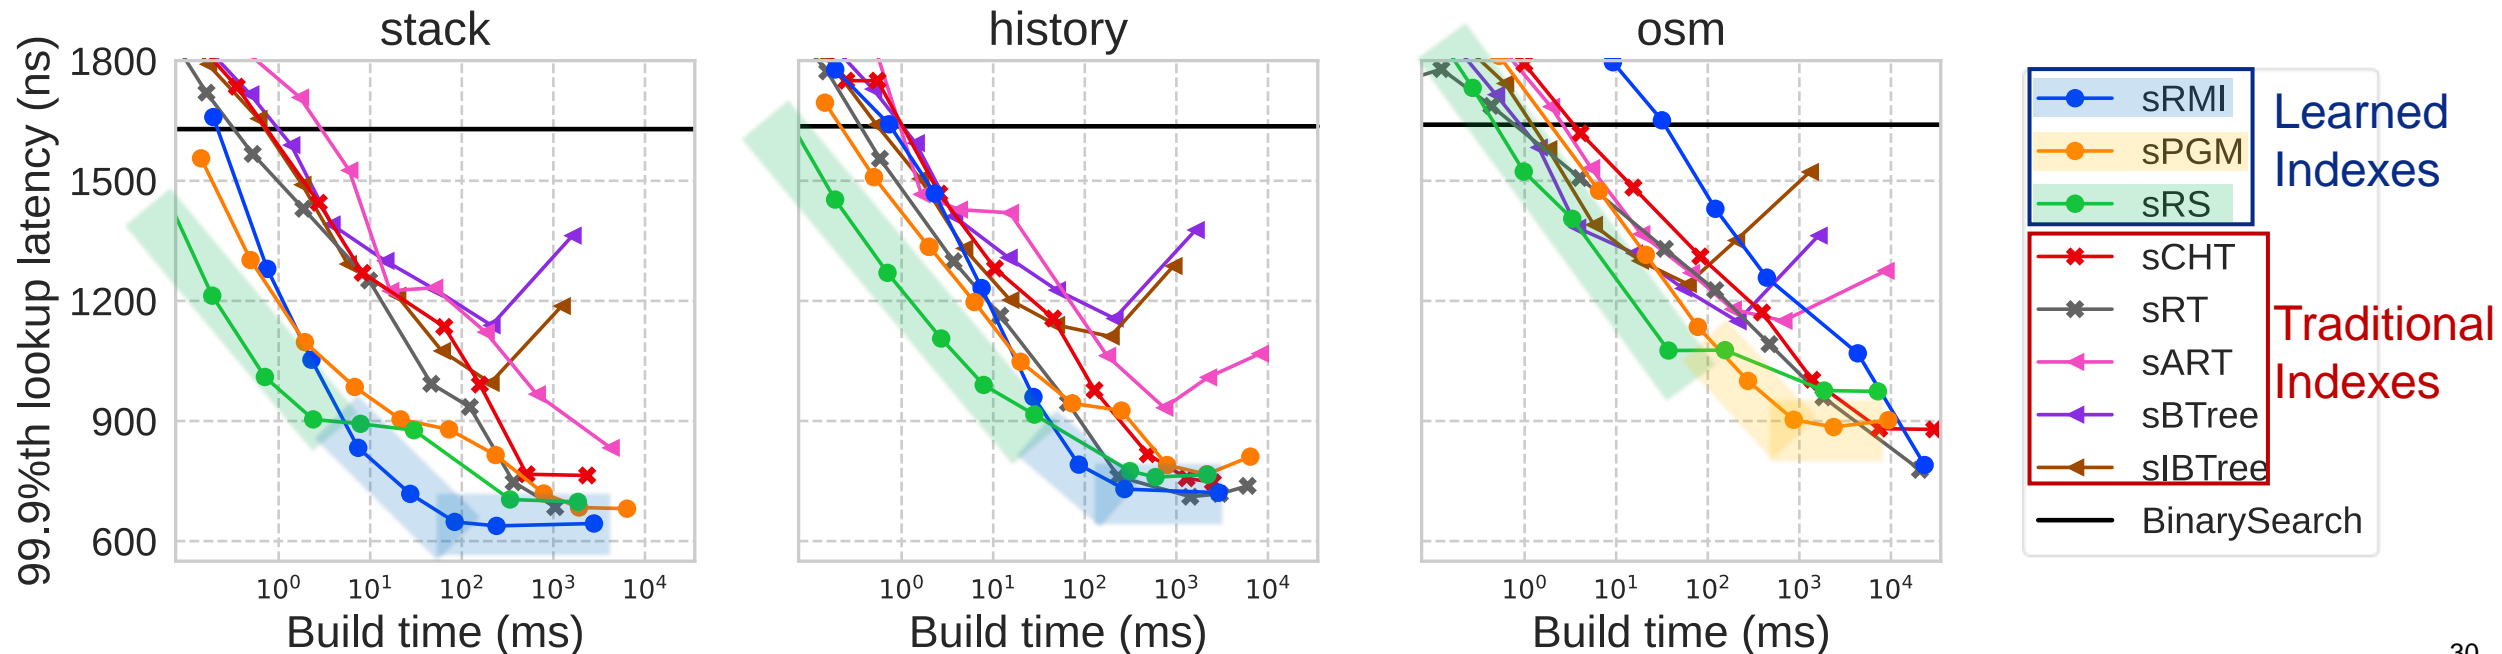
# 5. Evaluation

## ■ 4. Pareto Optimal Analysis

- Question. Can learned indexes be built efficiently (in terms of build time and lookup latency than traditional indexes through sampling)?

- **Pareto optimal** (build-efficient) in terms of build time and **tail** lookup latency

➤ no alternative that has both shorter build time and lower latency



## 6. Conclusion

1. Learned indexes are space-efficient, but long build time make it impractical.
2. Sampling has 3 challenges : 1) losing error-bound property, 2) absence of benchmark, and 3) complex sampling trade-offs.
3. We propose 1) novel sample learning algorithms which preserves error-bound , 2) new benchmark 🌿 BASIL, and 3) analysis of sampling trade-offs.
4. We show that sampling can 1) expand the design space, 2) reduce build time without significant performance loss, and 3) build learned indexes efficiently.

**Thank you**