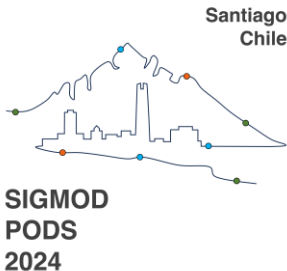


Can Learned Indexes be Built Efficiently?

A Deep Dive into Sampling Trade-offs

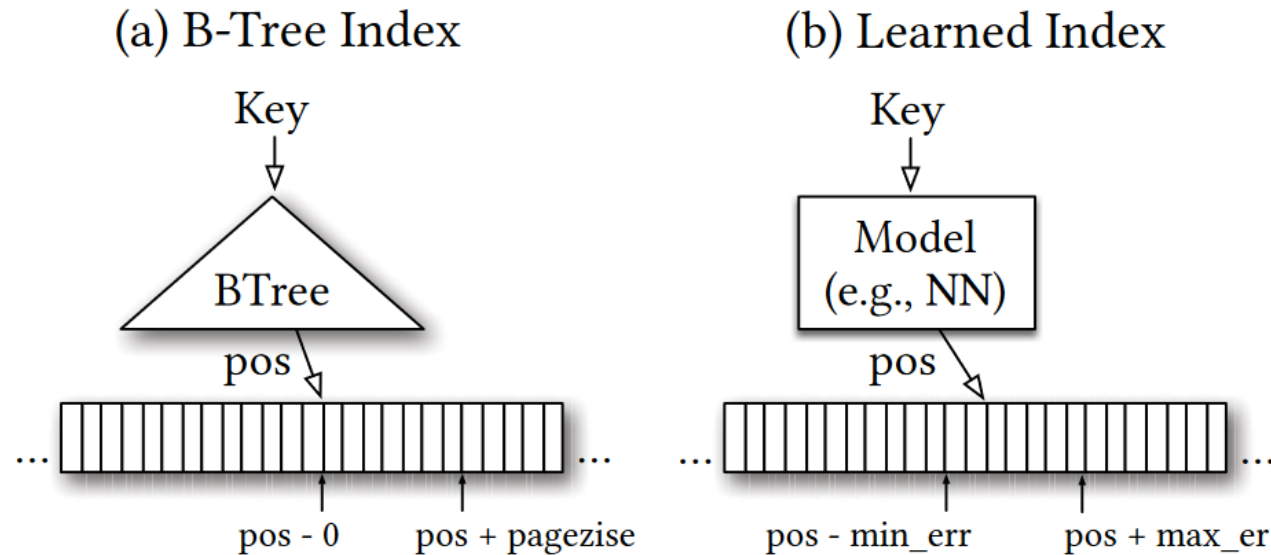
Minguk Choi, Seehwan Yoo, Jongmoo Choi

Dankook University, South Korea



1. Introduction

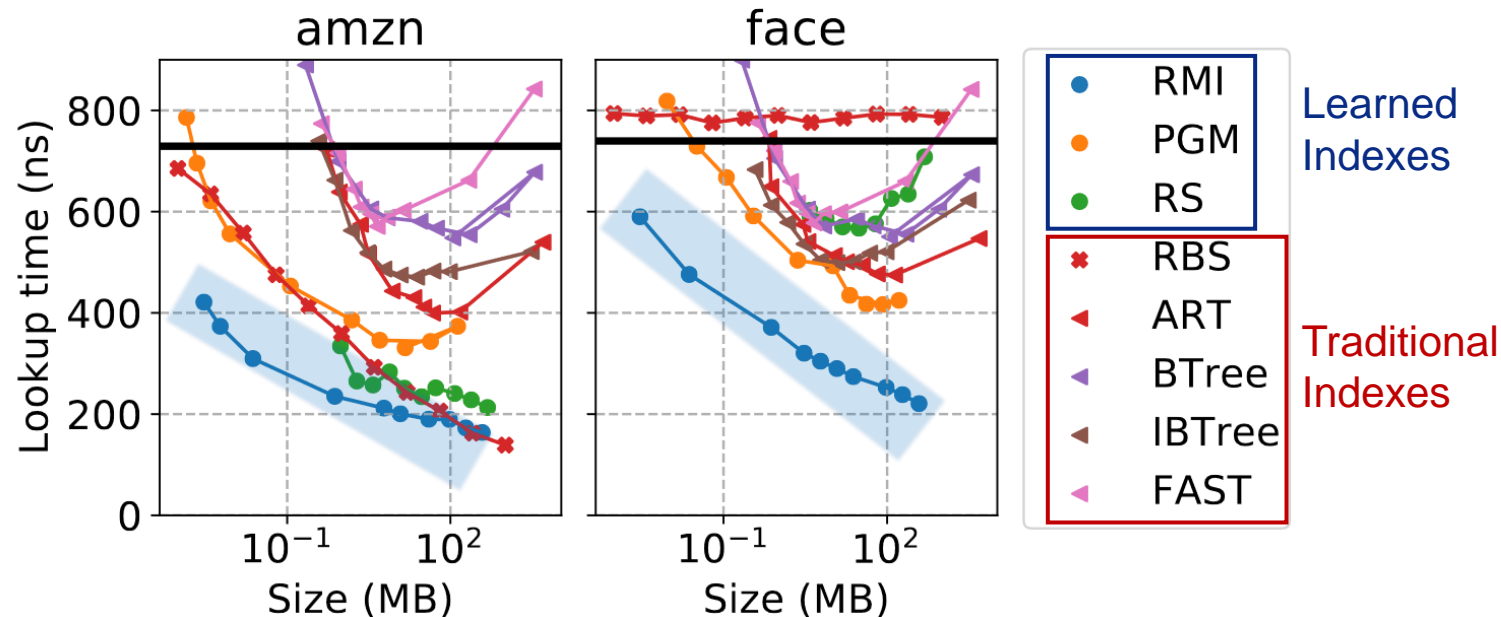
- Learned index
 - Index structure employs machine learning techniques
 - View the index as a model that **predicts the position of a key** in sorted array



(The Case for Learned Index Structures, SIGMOD '18)

1. Introduction


- Learned index
 - **Space-efficient** by effectively compressing data distribution through model
 - Pareto optimal in terms of index size and lookup latency in read-only workload
 - No alternative exists that has both a smaller size and lower latency



(Benchmarking Learned Indexes, VLDB '20)

1. Introduction


- Limitation of learned index : **Long index build time**
 - Significantly (up to about 2,000x) longer than traditional indexes



1. Introduction

■ Limitation of learned index : **Long index build time**

- Significantly (up to about 2,000x) longer than traditional indexes
 - Not Pareto optimal (build-efficient) for build time and lookup latency
- Still, there are **application** (e.g., LSM-Tree) where the index build time is crucial




1. Introduction

- Limitation of learned index : **Long index build time**
 - Significantly (up to 1,950x) longer than traditional indexes



Long build time has been identified as **a high priority** for future work in various papers

RMI (SIGMOD `18), RadixSpline (aiDM `20), PGM-Index (VLDB `20), SOSD (VLDB `20), Critical-RMI (VLDB `22)



1. Introduction

- The primary reason for long build time of learned index

$$\text{Index build time} = \text{Per - element overhead} \times \text{Number of elements}$$

- 1) Higher per-element overhead
- 2) Complete traversal and training

- To mitigate per-element overhead
 - Light-weight model : RadixSpline (aiDM `20), Bourbon (OSDI `20)
 - It still shows longer build time than traditional indexes

1. Introduction

- Our Approach : **Sampling**


- While sampling may seem simple and even naïve, it is indeed complex

- Challenges

1. **Absence** of benchmark for sampling applied indexes
2. **Losing** error-bound property due to sampling loss
3. **Complex** trade-offs in terms of model, index, and micro-architecture


2. Background

- Workload : Read-only in-memory
 - Practical beginning point of learned index
 - Dataset (D) : Sorted array of unique integer keys without duplicates
 - Lookup : Find the position of a lookup key k in D
 - ① Prediction : Estimate the position of k as p
 - ② Correction : Find the exact position of k based on p




2. Background

- Workload : Read-only in-memory
 - Error-bound property
 - $\forall k \in D, Error(k) = |Pred(k) - Pos(k)| \leq \varepsilon$ ($= error - bound$)
 - k exist in correction range ($= [p - \varepsilon, p + \varepsilon]$) \rightarrow binary search
 - Important for robustness,
 - Especially where correction is expensive
 - E.g., Disk or remote I/O environments



3. Design



1. Unified sampling algorithm & implementation

 **BASIL** (Benchmark of Sampling Applied Learned Indexes)

1) Unified sampling algorithm

- **Systematic** sampling : extract every I^{th} (I = sampling interval) key from 0-th key to the last key
 - ✓ Pros : Simple, universal, no decision/reordering cost
 - ✓ Cons: Not optimal (other methods, e.g., adaptive, should be explored)

2) Unified sampling implementation


- All indexes access and train only sample key-value data from entire dataset

3. Design

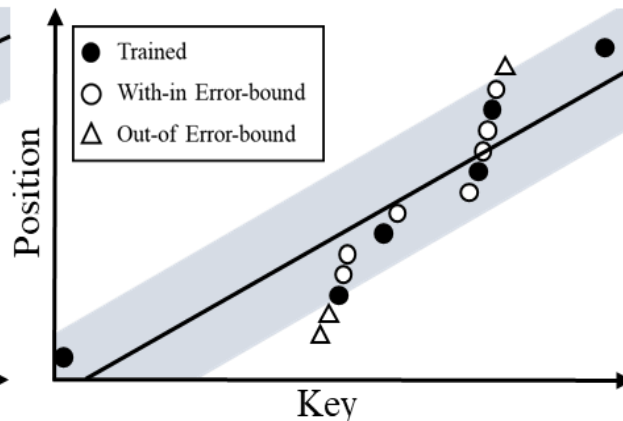
2. Sample learning algorithm

- EB-PLA (Error-bounded Piece-wise Linear Approximation) model
 - Without sampling, model trains **all** keys with error-bound $\varepsilon \rightarrow \text{Error}(k) \leq \varepsilon$
 - With sampling, model trains only **sample** I^{th} keys with error-bound $\varepsilon \rightarrow \text{Error}(k) \leq \varepsilon$
 - **Losing** of the error-bound property, which is learning objective of the model

(a) Train All ($I = 1$)
with $\varepsilon = 3$



(b) Train Sample ($I = 3$)
with $\varepsilon = 3$




3. Design

2. Sample learning algorithm


- Sample EB-PLA algorithm

- Refine error-bound due to sampling loss ($\varepsilon' = \varepsilon + I - 1$)
 - In Fig. (c), preserves error-bound property, but cannot guarantee desired error-bound (ε)
- Replace the learning error-bound to $\delta (= \varepsilon - I + 1)$ for desired error-bound (ε)
 - In Fig. (d), preserves error-bound (ε) by learning less data with smaller & stricter error bound


(a) Train All ($I = 1$)
with $\varepsilon = 3$




(b) Train Sample ($I = 3$)
with $\varepsilon = 3$



(c) Refine Error-bound
 ε from 3 to 5



(d) Train Sample ($I = 3$)
with $\delta = 1$

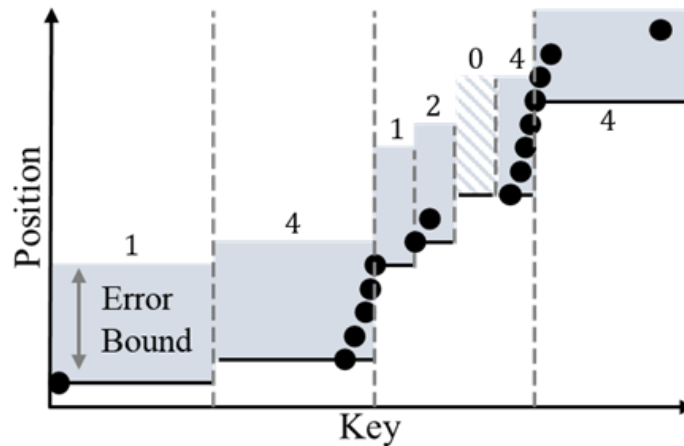


3. Design

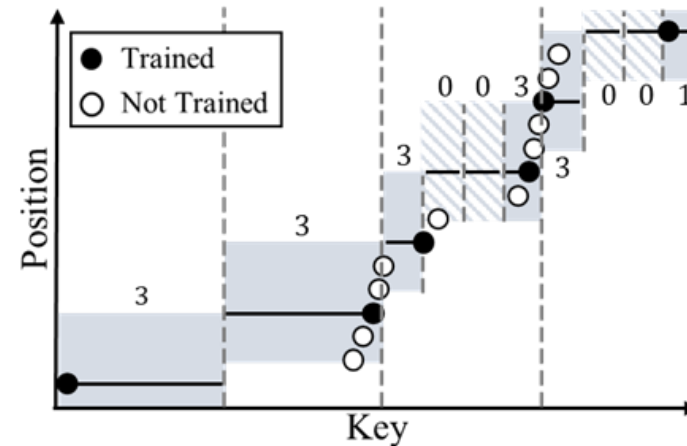
2. Sample learning algorithm

- Sample EB-Histogram algorithm

- Fig (a), Train all with error-bound $\varepsilon \rightarrow \forall k \in D, k \in [p, p + \varepsilon]$
- Fig (b), Train sample with error-bound $\delta (= \varepsilon - I + 1) \rightarrow \forall k \in D, k \in [p - I + 1, p + \delta]$
 - Preserve Correction length ($\varepsilon + 1 = \delta + I$)



(a) EB-Histogram ($I = 1, \varepsilon = 5$)

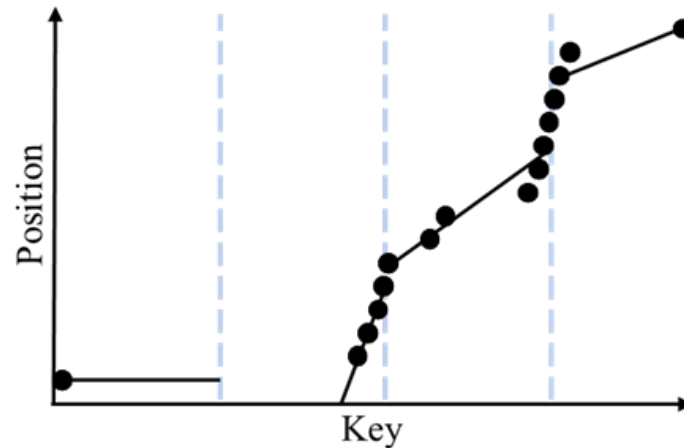


(b) Sample EB-Histogram ($I = 1, \delta = 3, \varepsilon = 5$)


3. Design

2. Sample learning algorithm

- PLR with Simple Linear Regression
 - Fig. (a), Model itself cannot guarantee an error-bound regardless of sampling
 - To guarantee error-bound, Measuring all data errors after training causes significant overhead.
 - Fig. (b), Sampling can decrease accuracy(MSE), but error-bound property doesn't change.



(a) Simple Linear Regression ($I = 1$)



(b) Simple Linear Regression ($I = 3$)

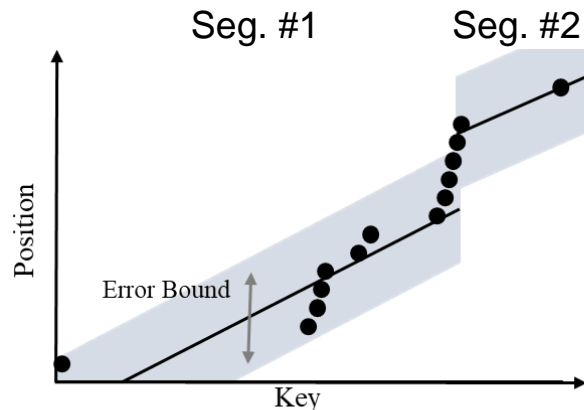
3. Design

3. Internal changes due to sampling

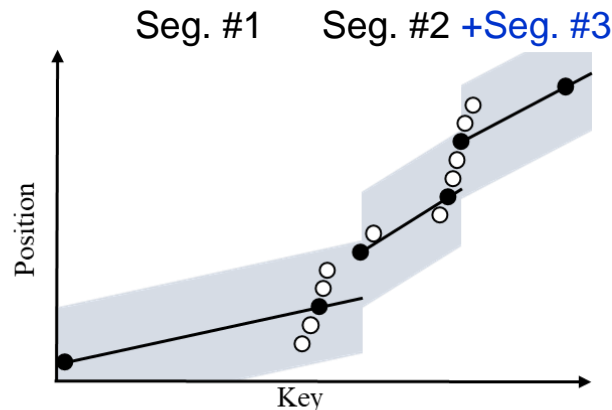
- Depend on segmentation manner

1) **Dynamic** segmentation (EB-PLA, EB-Histogram)

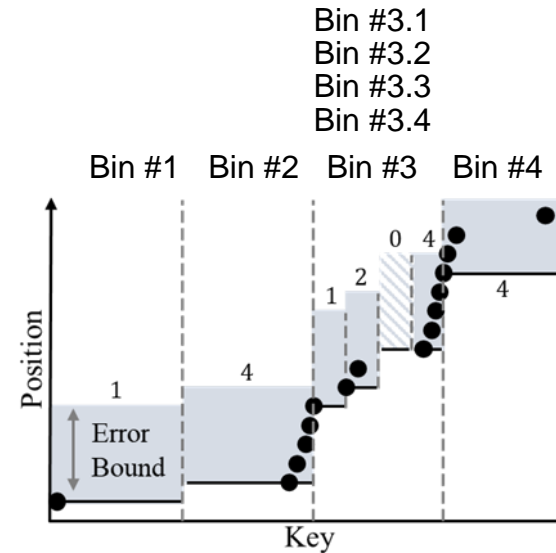
- Definition : Dynamically segment key ranges according to the distribution
- Trade-off : Decrease build time but aggressive sampling can **increase # of segments (bins)**



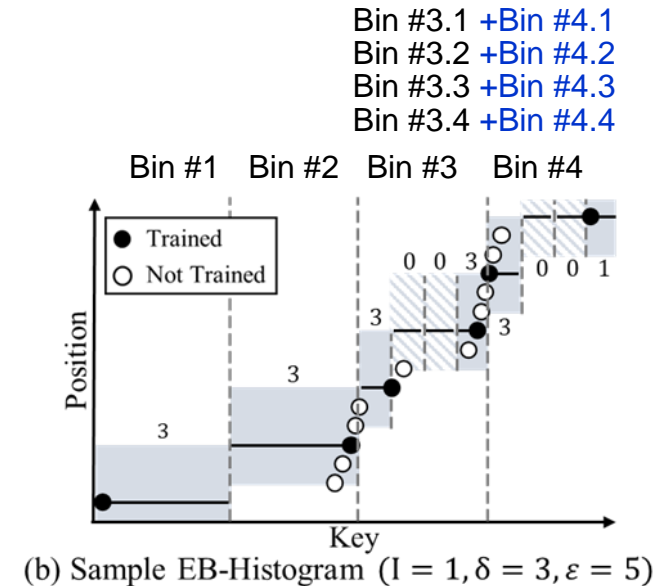
(a) EB-PLA ($I = 1, \varepsilon = 3$)



(b) Sample EB-PLA ($I = 1, \delta = 1, \varepsilon = 3$)



(a) EB-Histogram ($I = 1, \varepsilon = 5$)



(b) Sample EB-Histogram ($I = 1, \delta = 3, \varepsilon = 5$)

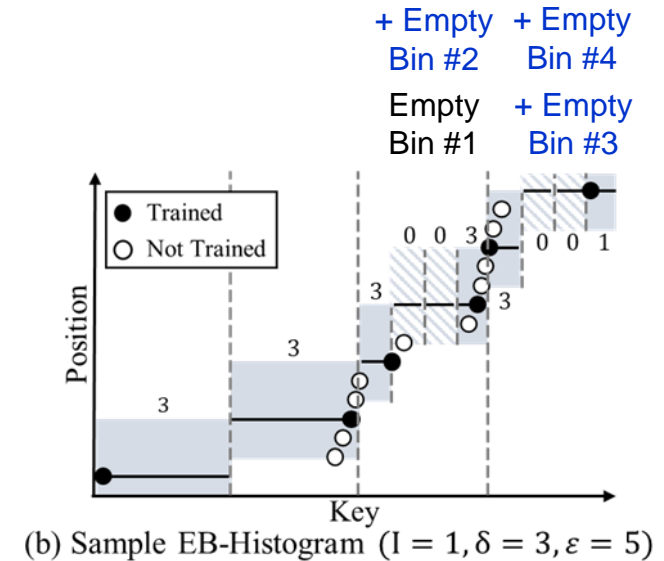
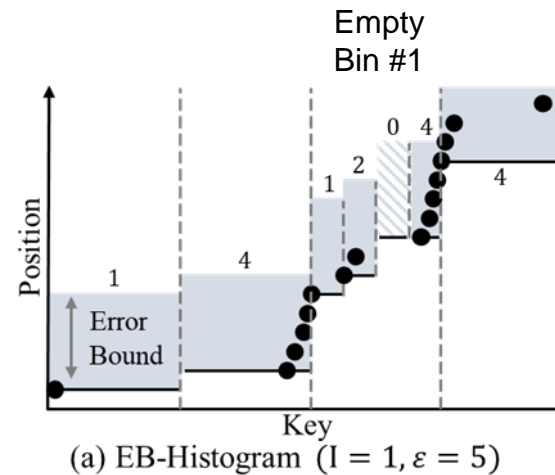
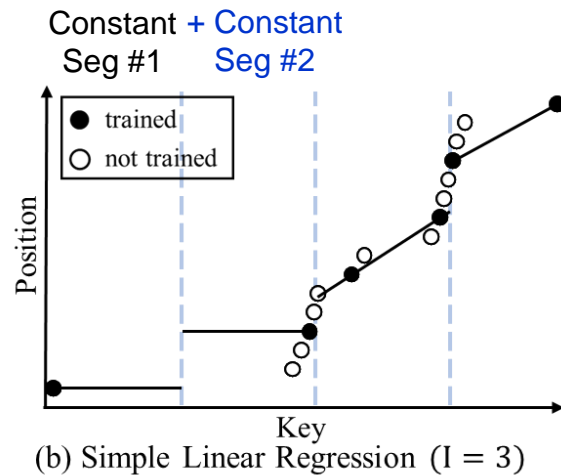
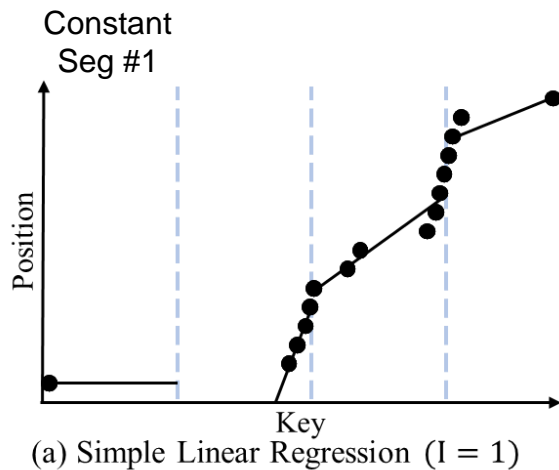
3. Design

3. Internal changes due to sampling

- Depend on segmentation manner

2) Fixed segmentation (PLR, EB-Histogram)

- Definition : Segment key ranges into a fixed number of segments
- Trade-off : Decrease build time but aggressive sampling can increase # of underfitting segments



4. Evaluation Setup



BASIL (Benchmark of Sampling Applied Learned Indexes)

- Applied sampling to 7 indexes, prefixed with “s”
 - 3 Learned, 2 Histogram, 3 Tree-based indexes

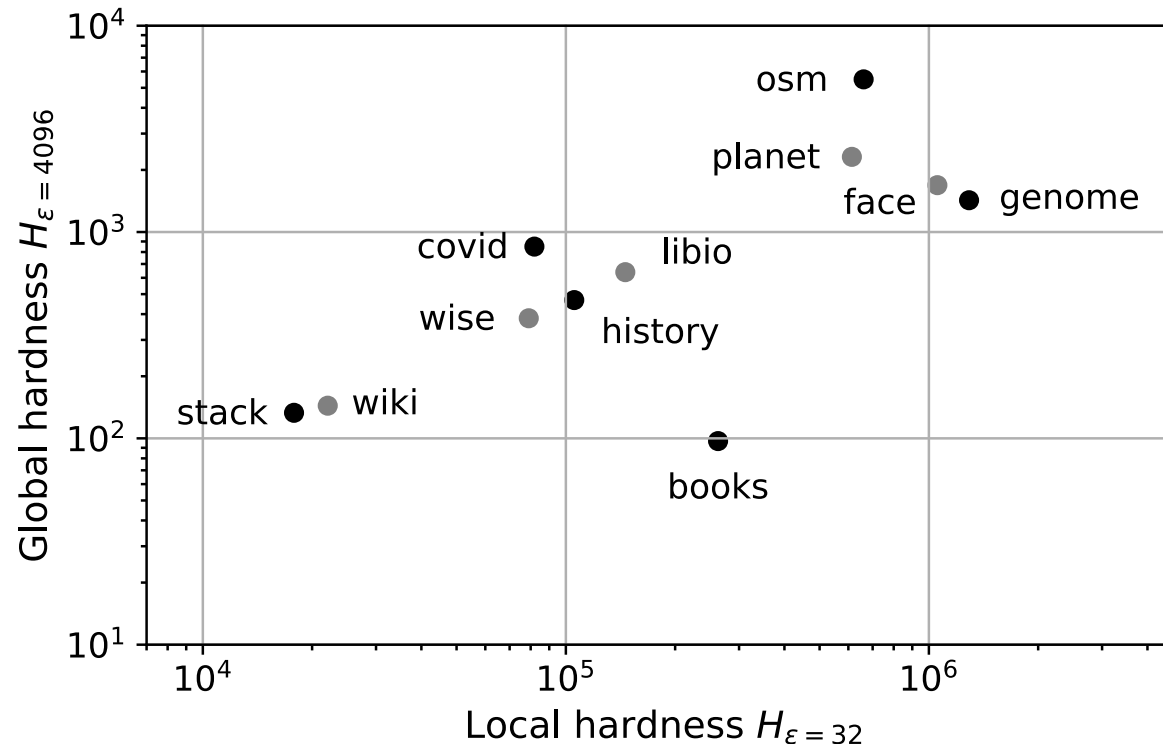
| Type | Index | Internal Model | Correction Search |
|------------|---------------------------|-----------------------------------|--------------------|
| Learned | sRMI | Simple Linear Regression | Exponential Search |
| Learned | sPGM / sRS | Sample EB-PLA | Binary Search |
| Histogram | sCHT | Sample EB-Histogram (Equal-width) | |
| Histogram | sRT | Sample Histogram (Equal-width) | |
| Tree-based | sART / sB+-Tree/ sIB-Tree | - | |

4. Evaluation Setup



BASIL (Benchmark of Sampling Applied Learned Indexes)

- Datasets : 6 representative datasets with 200 million key-value pairs
- Workload : Lookup uniform random 10 million keys from the dataset.
- Environment: Intel(R) Xeon(R) Gold 6338 CPU 2.00 GHz, 48 MB L3 with 512 GB of main memory



5. Evaluation

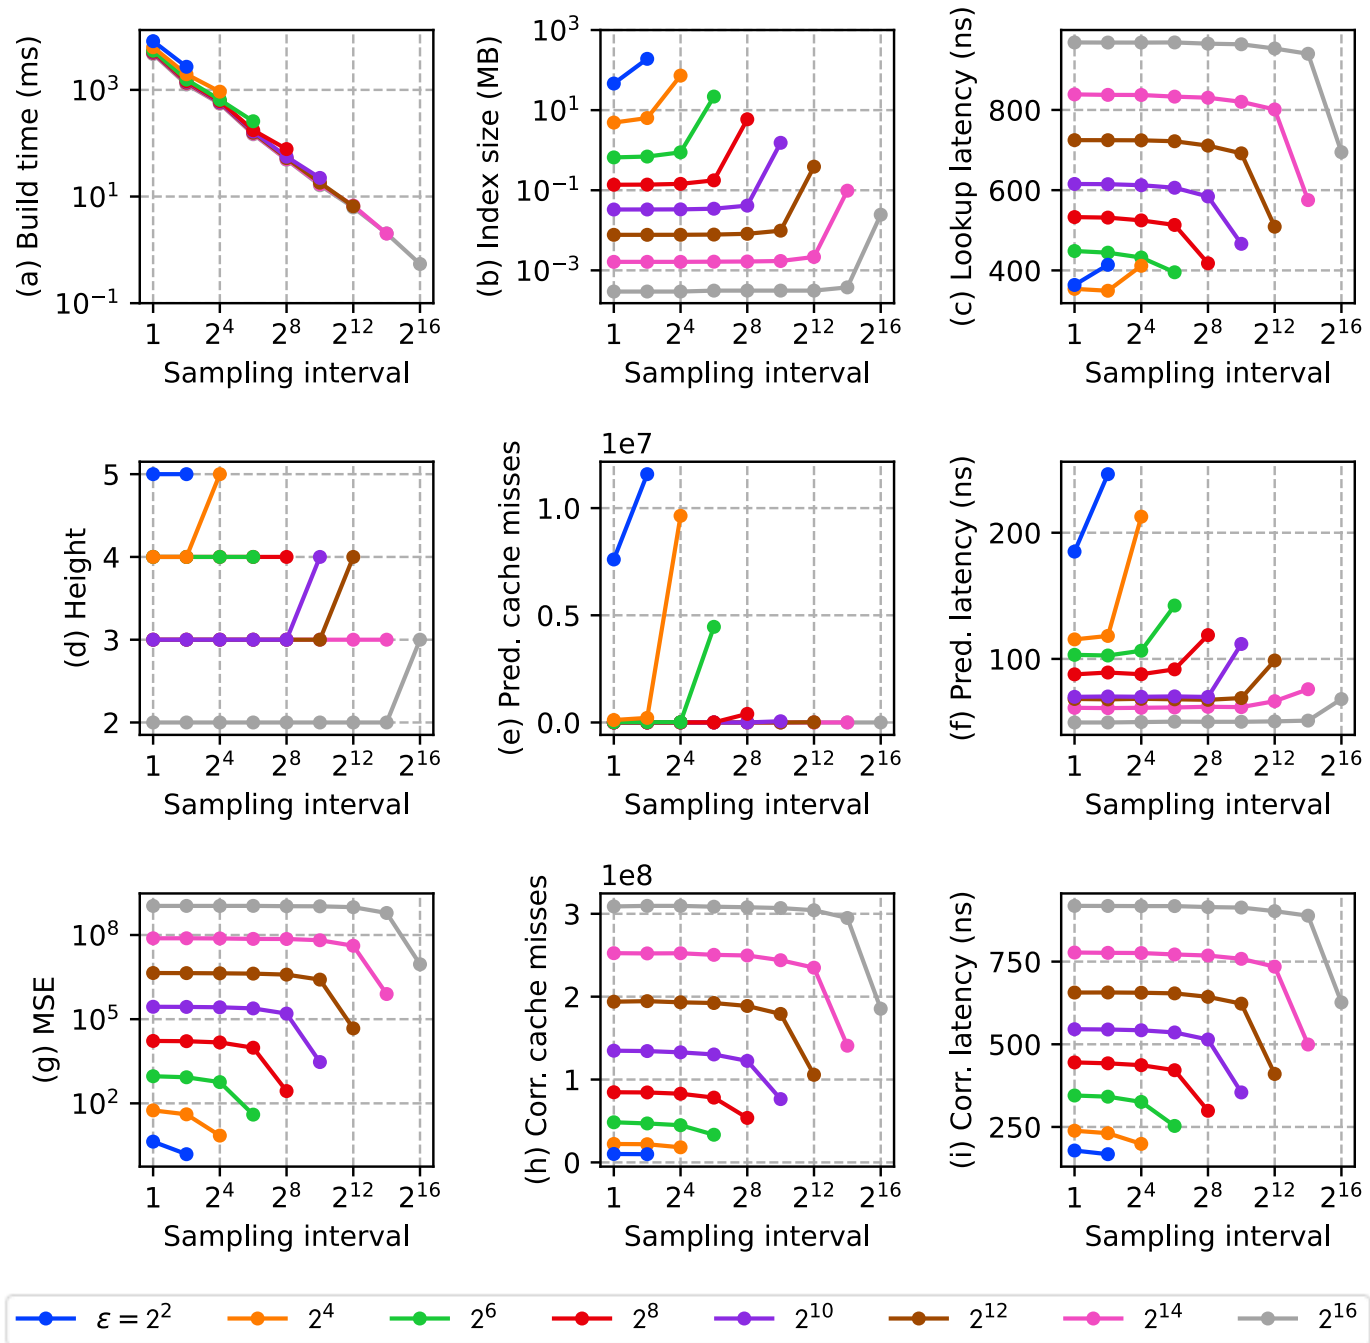
1. Sampling Trade-offs

■ Index

- sPGM with Sample EB-PLA

■ Metrics

- Index:** (a) Build Time, (b) Size, (c) Latency, (f) Pred. latency, (i) Corr. latency
- Model:** (d) Height, (g) MSE (Accuracy)
- Micro-architecture:** (e) Pred. Cache Miss, (f) Corr. Cache Miss

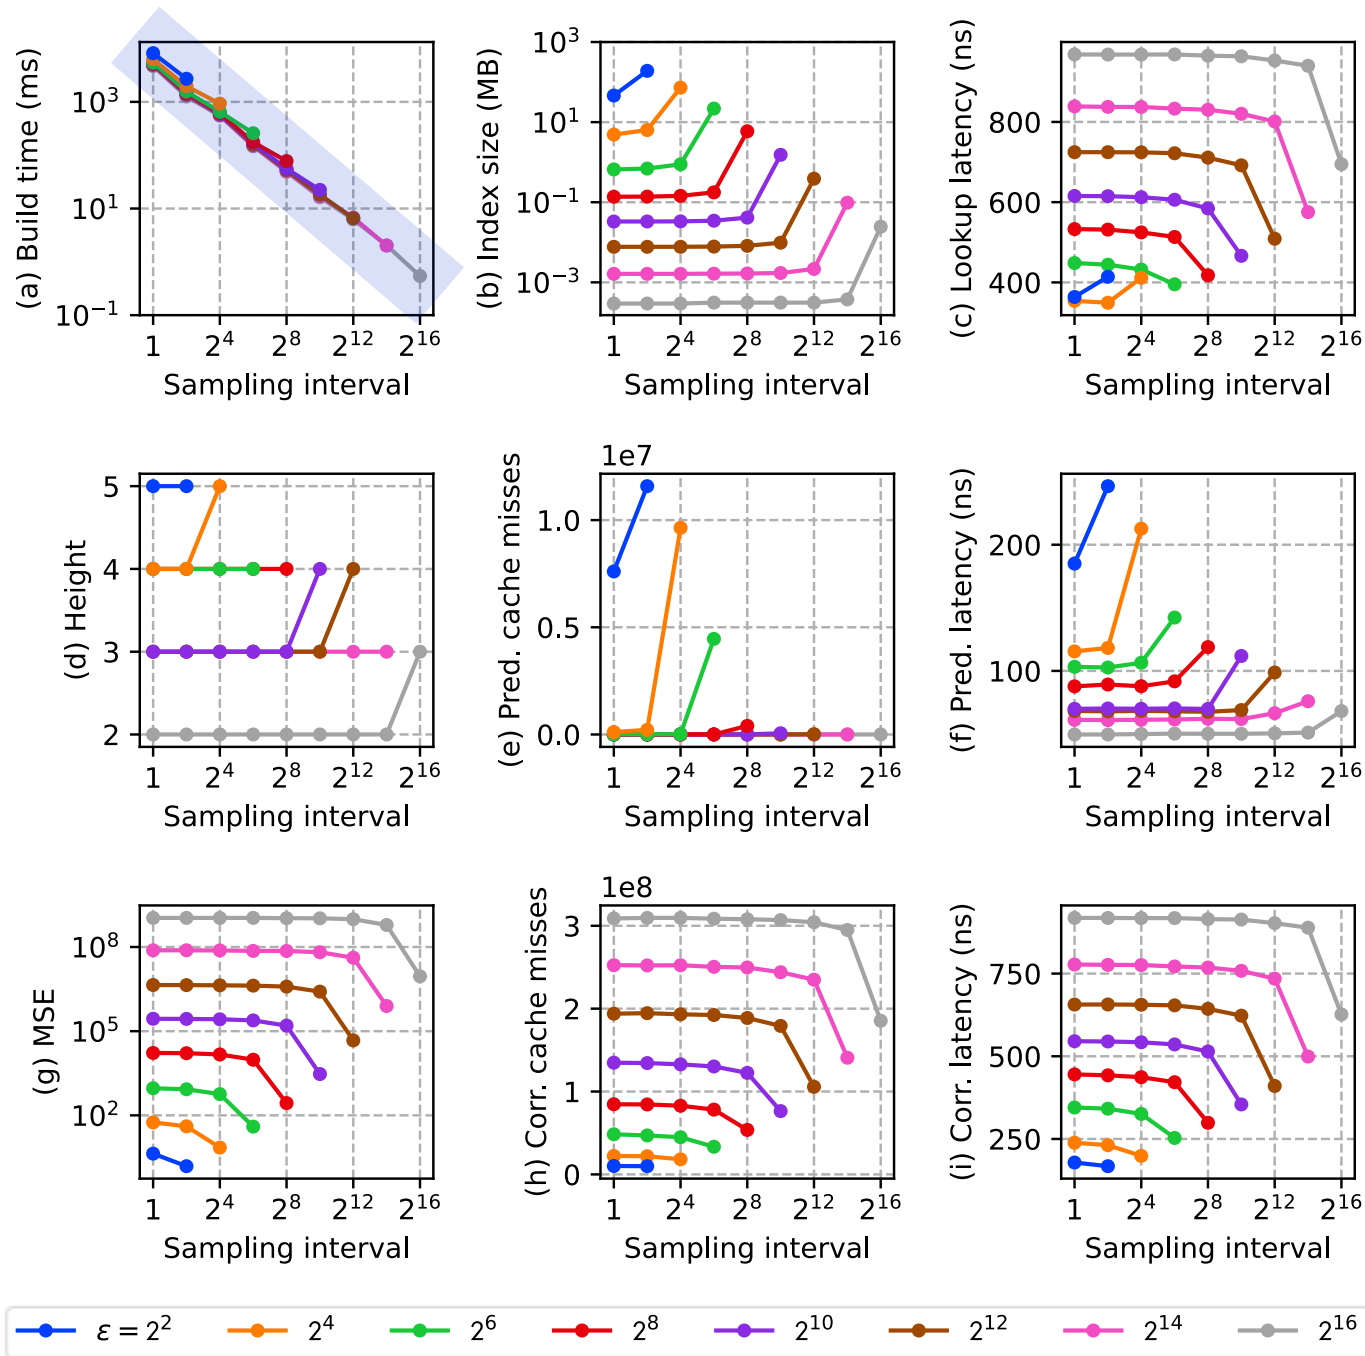


5. Evaluation

1. Sampling Trade-offs

- When sampling interval (I) increases, **(a) build time decreases** by order of magnitude

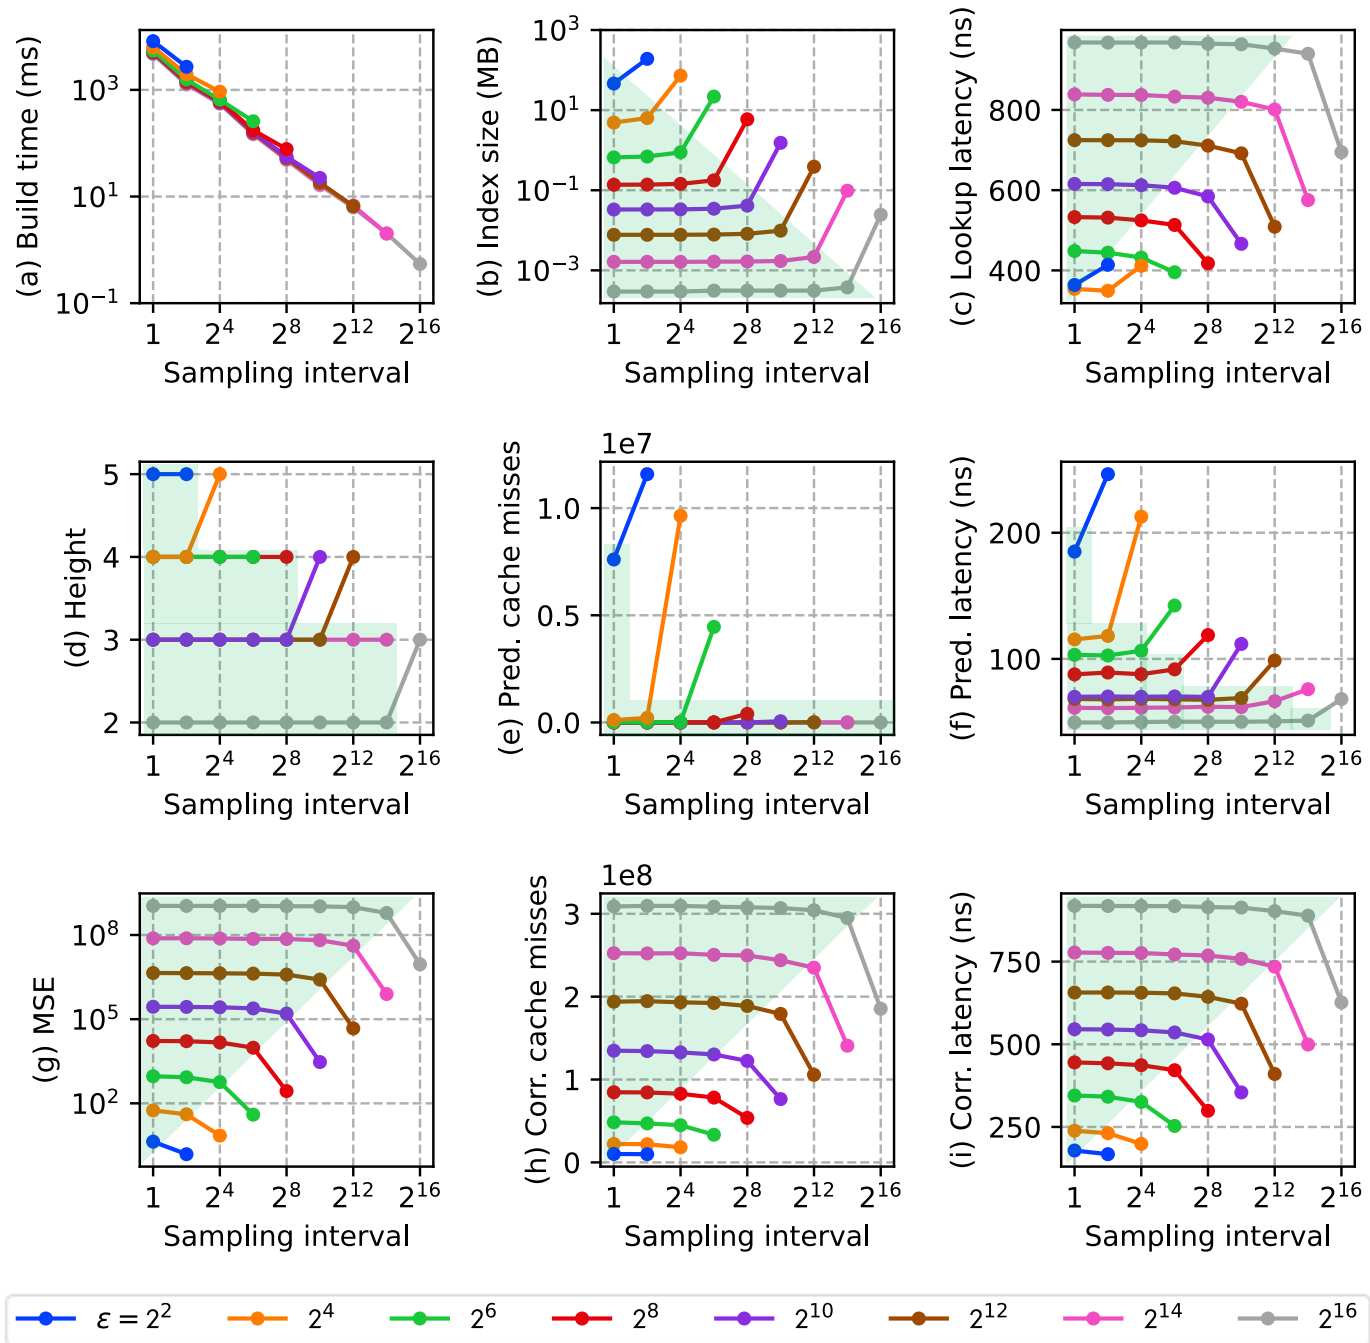
Dataset : History, Error bound ($\epsilon \in [2^2, 2^{16}]$), Sampling interval ($I \in [2^0, \epsilon (\leq 2^{16})]$)



5. Evaluation

1. Sampling Trade-offs

- Each error-bound has **threshold interval (I_{TH})**
 - mostly $\varepsilon = I^{TH}$
- Until** threshold interval (I_{TH}), (b-i) the rest of metrics **remain consistent**



5. Evaluation

1. Sampling Trade-offs

- **After** threshold interval I_{TH} ,

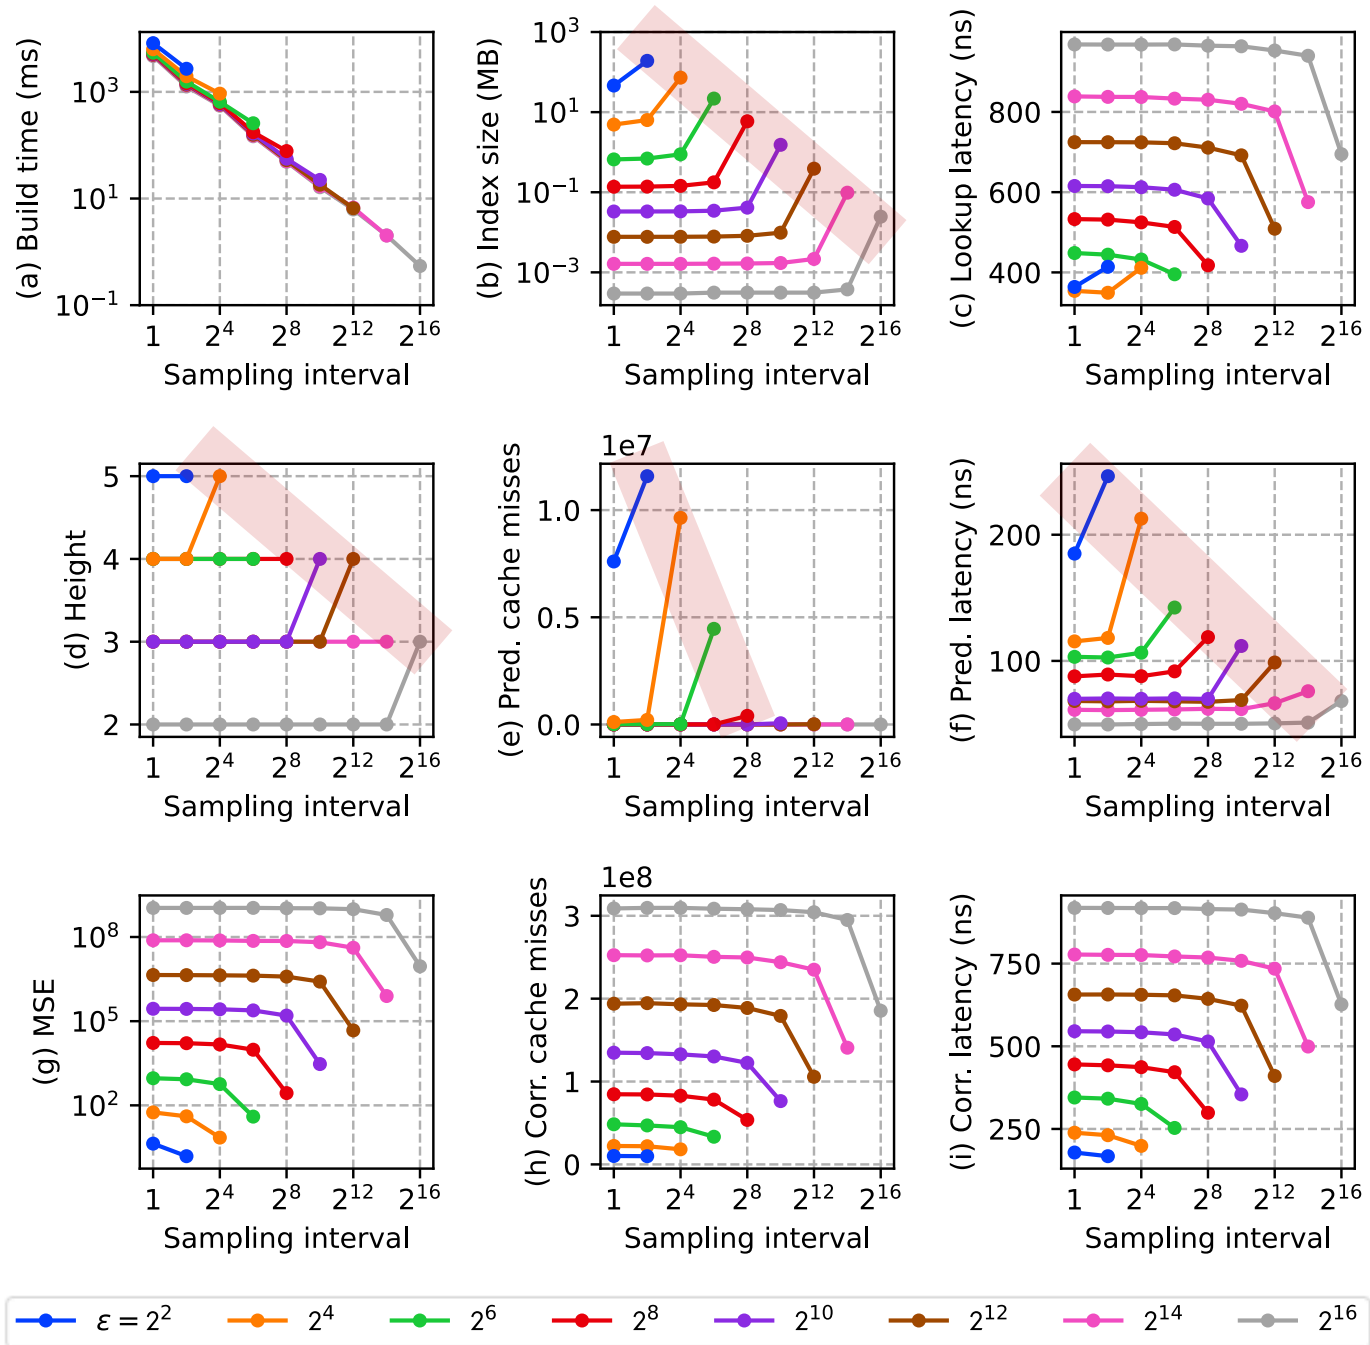
of linear segments \uparrow

→ (b) Size \uparrow

(d) Height \uparrow

→ (e) Pred. cache miss \uparrow ,

(f) Pred. latency \uparrow



5. Evaluation

1. Sampling Trade-offs

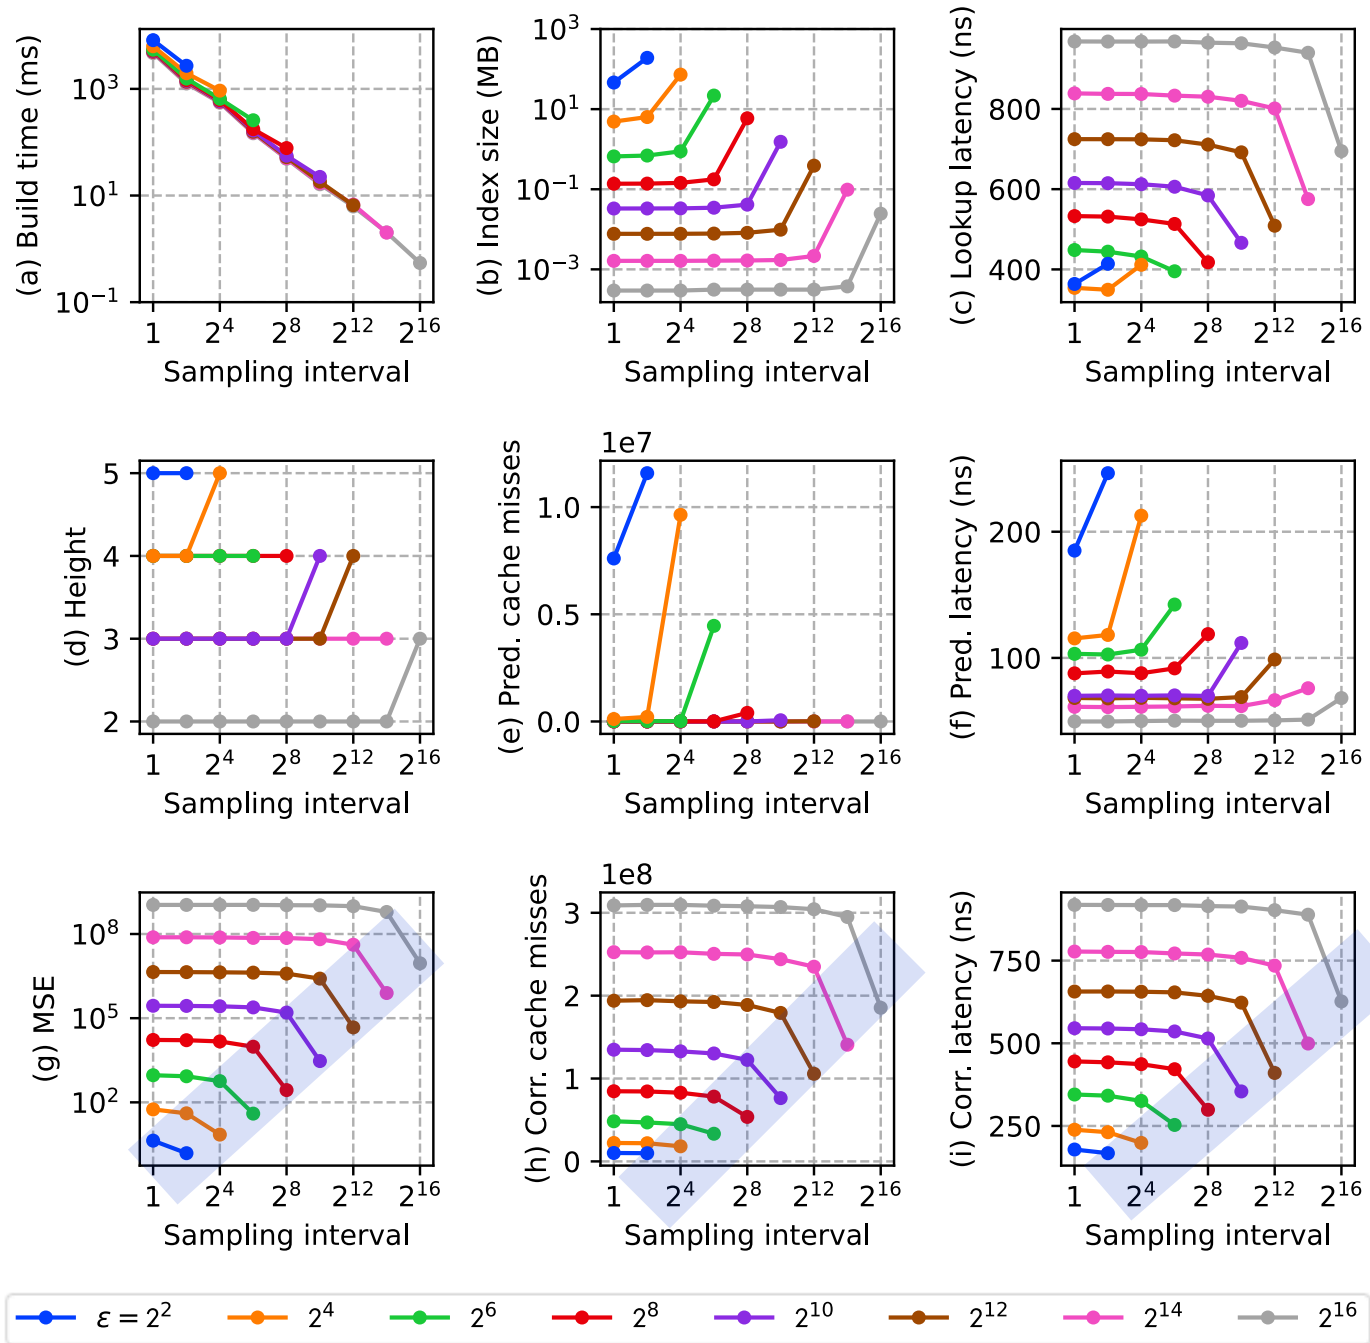
- After threshold interval I_{TH} ,

of linear segments \uparrow

→ (g) MSE \downarrow

→ (h) Corr. cache miss \downarrow

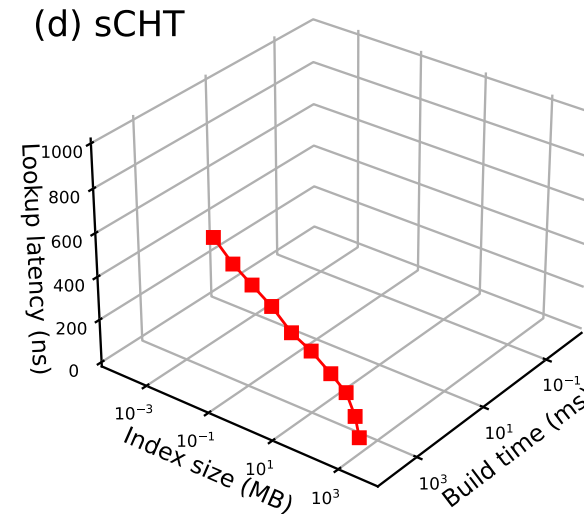
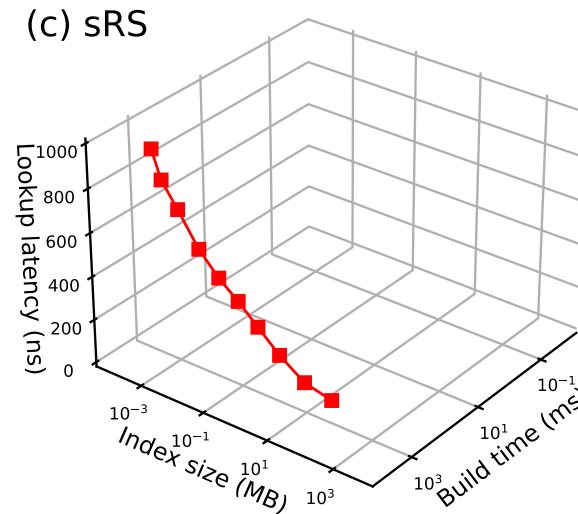
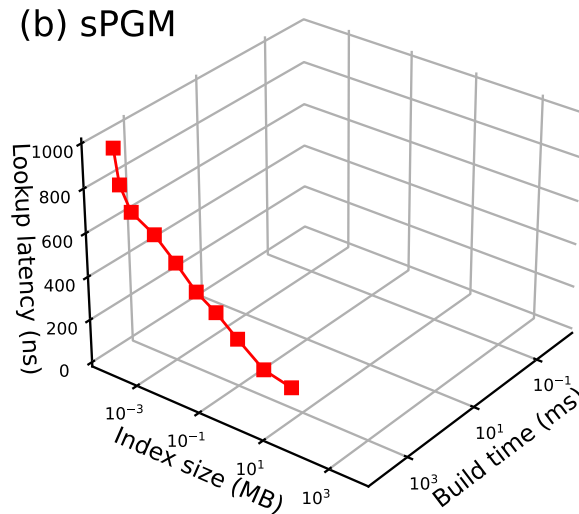
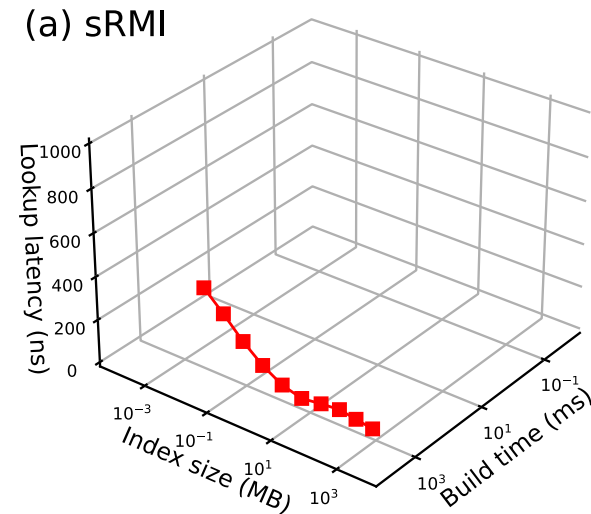
(i) Corr. latency \downarrow



5. Evaluation

■ 2. Design Space of Learned Indexes

- **Absence** of trade-offs between build time, index size, and lookup latency
 - Incur significant build times regardless of size and lookup

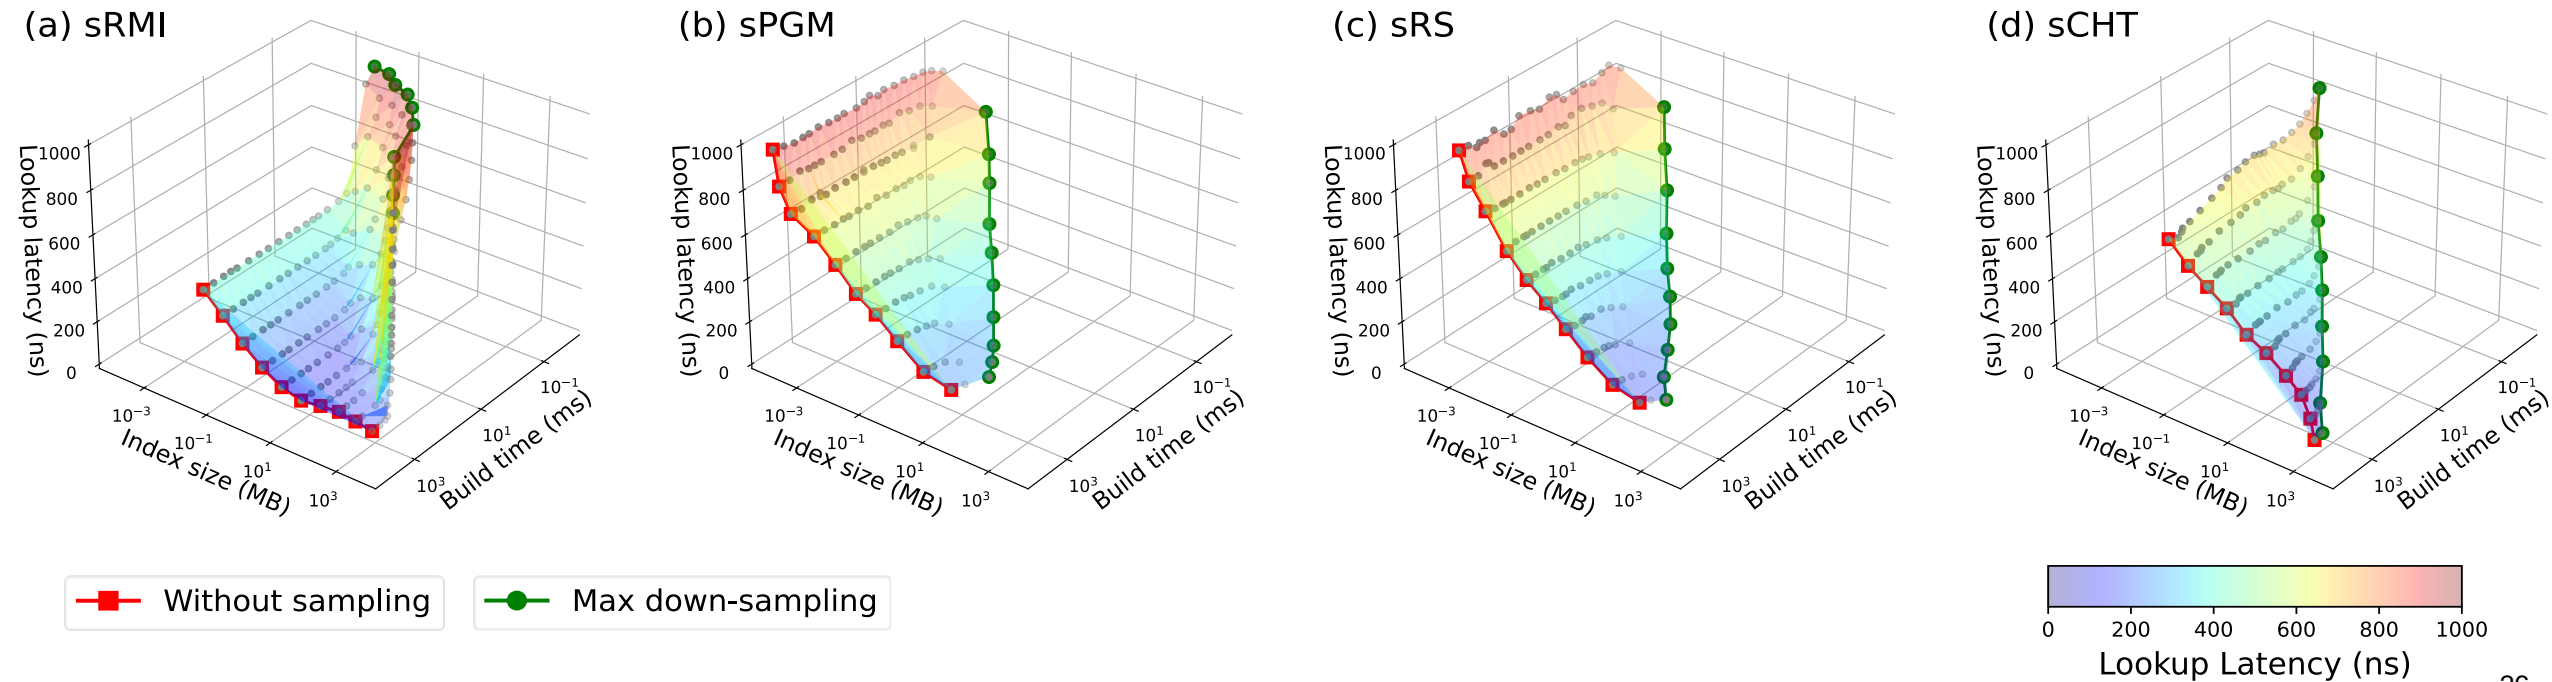


—■— Without sampling

5. Evaluation

2. Design Space of Learned Indexes

- Sampling **introduces** trade-offs between build-time, size, and lookup latency
 - Broaden** design space of learned indexes from 2D to 3D

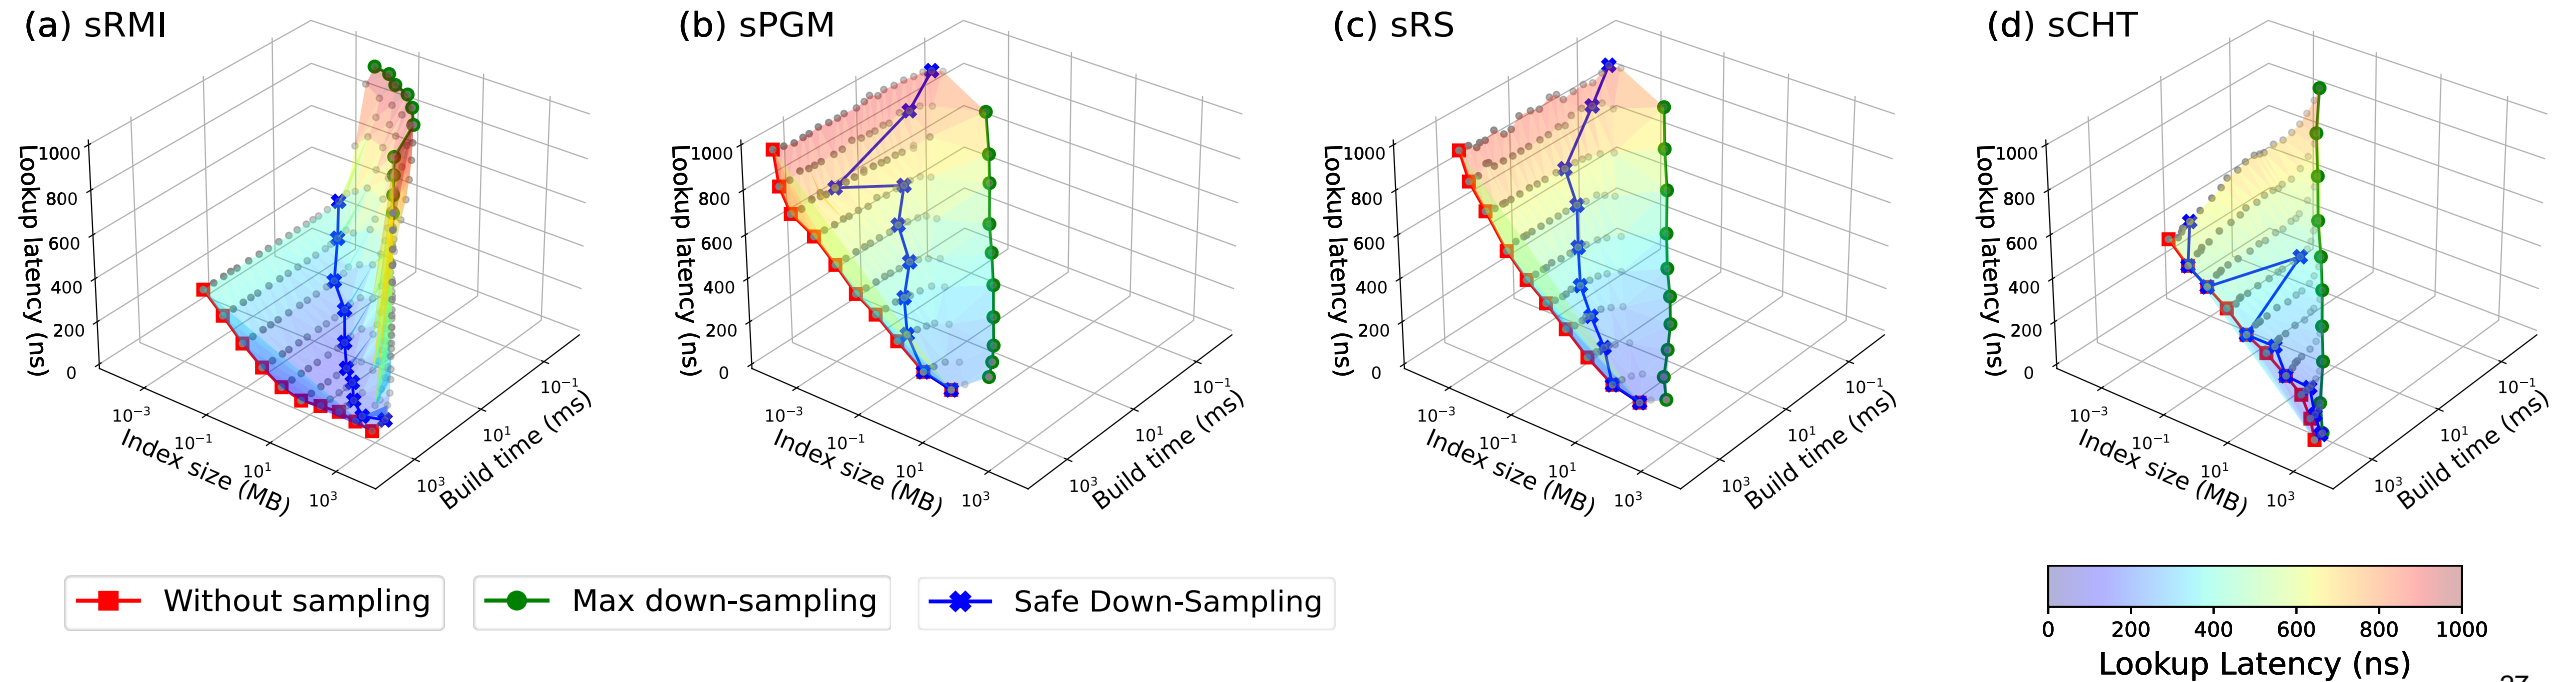


5. Evaluation

■ 3. Build Speed-up

- Question. How much can sampling reduce build time without significantly degrading index performance?

➤ **Safe** Down-sampling where size & lookup latency increase by **less than 5%**




5. Evaluation

3. Build Speed-up

- Question. How much can sampling reduce build time without significantly degrading index performance?

➤ **Safe** Down-sampling where size & lookup latency increase by **less than 5%**




5. Evaluation

■ 4. Pareto Optimal Analysis

- Question. Can learned indexes be built efficiently (in terms of build time and lookup latency than traditional indexes through sampling)?

- **Pareto optimal** (build-efficient) in terms of build time and **average** lookup latency

➤ no alternative that has both shorter build time and lower latency



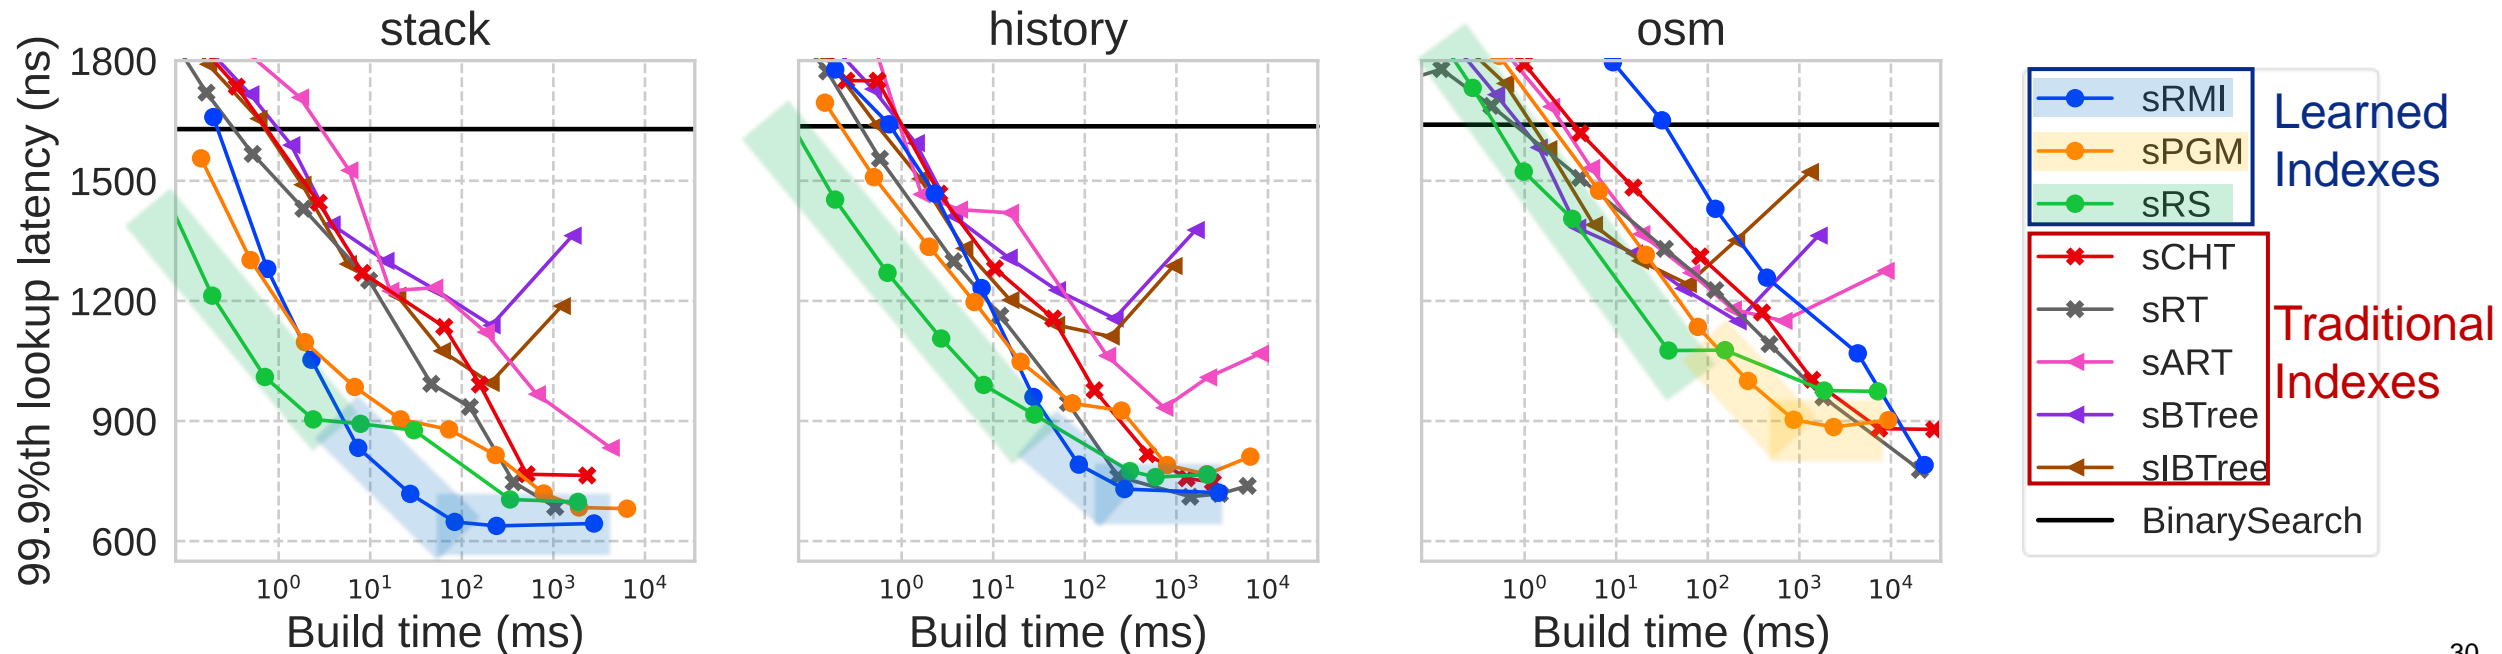
5. Evaluation

■ 4. Pareto Optimal Analysis

- Question. Can learned indexes be built efficiently (in terms of build time and lookup latency than traditional indexes through sampling)?

- **Pareto optimal** (build-efficient) in terms of build time and **tail** lookup latency

➤ no alternative that has both shorter build time and lower latency



6. Conclusion

1. Learned indexes are space-efficient, but long build time make it impractical.
2. Sampling has 3 challenges : 1) losing error-bound property, 2) absence of benchmark, and 3) complex sampling trade-offs.
3. We propose 1) novel sample learning algorithms which preserves error-bound , 2) new benchmark 🌿 BASIL, and 3) analysis of sampling trade-offs.
4. We show that sampling can 1) expand the design space, 2) reduce build time without significant performance loss, and 3) build learned indexes efficiently.

Thank you