

Minguk Choi

mgchoi@dankook.ac.kr | Homepage | LinkedIn | Github

RESEARCH INTEREST

Previous: Linux Kernel, Key-Value Store

Current: Learned Index Structures

Future: *ML for Systems* - Learned Data Structures, Learned Query Optimizer, Learned Cache
Systems for ML - Scheduler/Pipeline for ML-Cluster, Flash(CXL)-LLM

ON-GOING PROJECT

Semi-ordered Parallel Updatable Learned Index: Enhances performance in multi-threaded environments by reducing the critical section and data sorting overhead while maintaining lookup performance and size.

Adaptive Sample Training Algorithm for Learned Index: An algorithm for dynamically training with the optimal sampling interval by considering the data distribution and error-bound.

Accelerating Learned Index via SIMD: Accelerating the unique training and operations (e.g., model-biased insert) of learned index through the parallelism of internal ML models, which has not yet been actively studied.

EDUCATION

Dankook University
M.S. in AI-based Convergence

Yongin, South Korea
Mar 2023 – Aug 2024

Dankook University
B.S. in Software Science

Yongin, South Korea
Mar 2017 – Feb 2023

PUBLICATIONS

International Conference

(SIGMOD 2024) Can Learned Indexes be Built Efficiently? A Deep Dive into Sampling Trade-Offs
Minguk Choi, Seehwan Yoo, Jongmoo Choi

International Journal

(Electronics 2023) An Empirical Study of Segmented Linear Regression Search in LevelDB
Ramadhan Agung Rahmat, Minguk Choi, Yoojin Chung, Jongmoo Choi.

Domestic Conference

(KCC 2024) Analysis of RMI Using CPU-Optimized Search Algorithms (**Best Paper Award**)
Yejin Oh, Minguk Choi, Boseung Kim, Yongjie Zhu, Seehwan Yoo, Jongmoo Choi

(KCC 2024) Accelerating RMI Training with SIMD
Boseung Kim, Minguk Choi, Yeojin Oh, Yongjie Zhu, Seehwan Yoo, Jongmoo Choi

(KCC 2024) Performance Analysis of Batch Prediction Using SIMD in RMI
Yongjie Zhu, Minguk Choi, Yeojin Oh, Boseung Kim, Seehwan Yoo, Jongmoo Choi

(KCC 2024) Breakdown Internal Operations in Updatable Learned Index
Suhwan Shin, Minguk Choi, Nakyeong Kim, Seehwan Yoo, Jongmoo Choi

(KCC 2024) Analysis of Updatable Learned Indexes with Index Size Perspective
Nakyeong Kim, Minguk Choi, Suhwan Shin, Seehwan Yoo, Jongmoo Choi

(KSC 2022) Bloom Filter Optimization in LevelDB based on Hit-Ratio
Hansu Kim, Minguk Choi, Seehwan Yoo, Jongmoo Choi

(KSC 2022) Read performance analysis according to Compaction Trigger
Zhao Guangxun, Sangwoo Kang, Minguk Choi, Seehwan Yoo, Jongmoo Choi

(KSC 2022) LevelDB Cache Structure and Performance Analysis
Subin Hong, Minguk Choi, Seehwan Yoo, Jongmoo Choi

(KSC 2022) Per Key-Value Checksum Analysis on RocksDB
Suhwan Shin, Seyeon Park, Minguk Choi, Seehwan Yoo, Jongmoo Choi

EXPERIENCE

Student Researcher

System Software Laboratory

June 2021 – Present
Yongin, Korea

- Linux Kernel (CPU Scheduler, Memory Allocator, Block I/O Layer)
- Key-Value Storage (LevelDB, RocksDB)
- Learned Index (Read-only/Updatable, Sampled Learning, Semi-ordered, SIMD)

Teaching Assistant

Operating System Practice

Dankook University
Sep 2023

- Conduct 40 hours of lectures on system programming and operating systems
- Create a various practice (Sloppy Counter, Dining Philosophers Problem, Readers-Writers problem, myShell)
- Setting and grading final-exam questions

Operating Systems

Mar 2024 – Jul 2024

- Create an CPU scheduler simulator assignment (8-type scheduler, various workloads, auto-graded through gtest)
- Create a concurrent data structure assignment (Queue/BST, various workloads, auto-graded through gtest)
- Give a ext2 file forensics assignment

System Programming

Sep 2023 – Dec 2023

- Give a myls, myCreate, myCat, myCopy, and myShell assignments

Index Structure Study

Jan 2024 – Feb 2024

- Conduct an introductory lecture on (Learned) Index Structures
- Guidance for student research (5 domestic conference papers, 1 best paper award)

LevelDB Study

Jul 2022 – Aug 2022

- Conduct an introductory lecture on (Learned) Index Structures
- Hosting a LevelDB tuning contest using YCSB
- Guidance for student research (4 domestic conference papers, open-source analysis document)

Compulsory Military Service

Honorable Discharge

Jan 2019 – Oct 2020
Korean Air Force

OPENSOURCE CONTRIBUTION

BASIL

Benchmark of sampling applied learned indexes (To be released)

Owner

Leveldb Study/Wiki

Analysis document about LevelDB Key-Value Store

Owner

YCSB-CPP

Yahoo! Cloud Serving Benchmark(YCSB) written in C++

Contributer

Uftrace

Function graph tracer for C/C++/Rust/Python

Wiki Contributor

PGM-Index

State-of-the-art learned data structure

Bug Report

PROJECTS

Key-Value DB for Unstructured Big Data (SW StarLab)

IITP
June 2021 – Present

Key-Value Store with Predictable Latency Support

NRF
Mar 2022 - Present

Future Vehicle Security System

KEA
May 2023 - Present

HONORS AND AWARDS

Best Paper Award, Korea Computer Congress

LANGUAGE AND TECHNICAL SKILLS

Languages: Native korean, English (TOEIC 970/990, OPIc IH)

Programming Languages: c/c++, python

Developer Tools: gcc/g++, gdb, uftrace, git, matplotlib

PERSONAL INTERESTS

Fitness: Running, Weight lifting

Cooking: Korean Cuisine, Japanese Cuisine, Grilling

REFERENCES

Prof. Jongmoo Choi choijm@dankook.ac.kr

Software Science, Dankook University

Prof. Seehwan Yoo seehwan.yoo@dankook.ac.kr

Mobile Systems Engineering, Dankook University